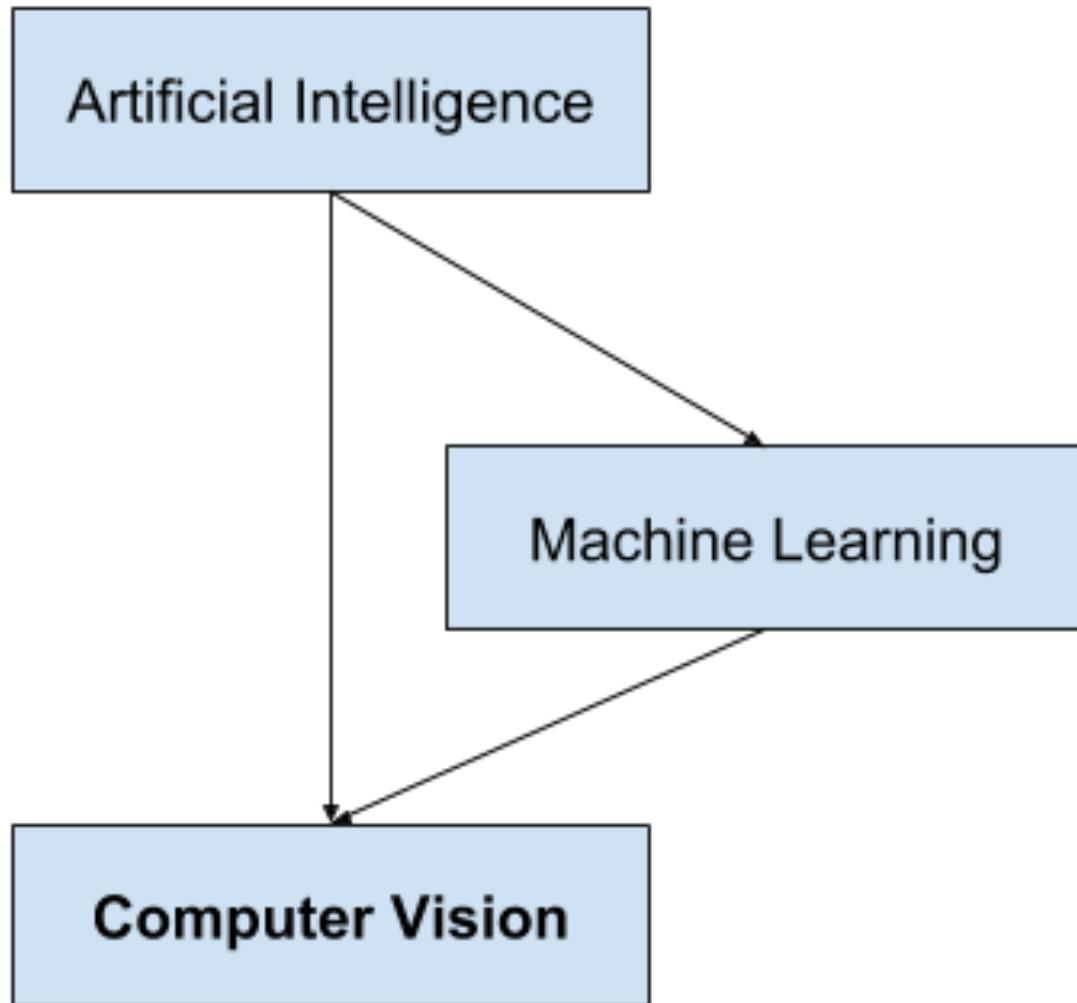


Computer Vision

Computer (Machine) vision is the automated extraction of information from images. At an abstract level, the goal of computer vision problems is to use the observed image data to infer something about the world. Information can mean anything from 3D models, camera position, object detection and recognition to grouping and searching image content.

Computer vision as a field is an intellectual frontier. Like any frontier, it is exciting and disorganized, and there is often no reliable authority to appeal to. Many useful ideas have no theoretical grounding, and some theories are useless in practice; developed areas are widely scattered, and often one looks completely inaccessible from the other.



Graphic vs Image

A graphic (image) is a representation of a 3D object. a digital representation of non-text information such as a drawing, chart, or photo.

An image (still Image, digital image) is the graphic converted for 2d display. the binary representation of any type of visual information like drawings, individual video frames, logos, pictures, graphs etc.

c.f) Rendering - an operation that takes graphics data and produces image data.

GRAPHIC / VECTOR

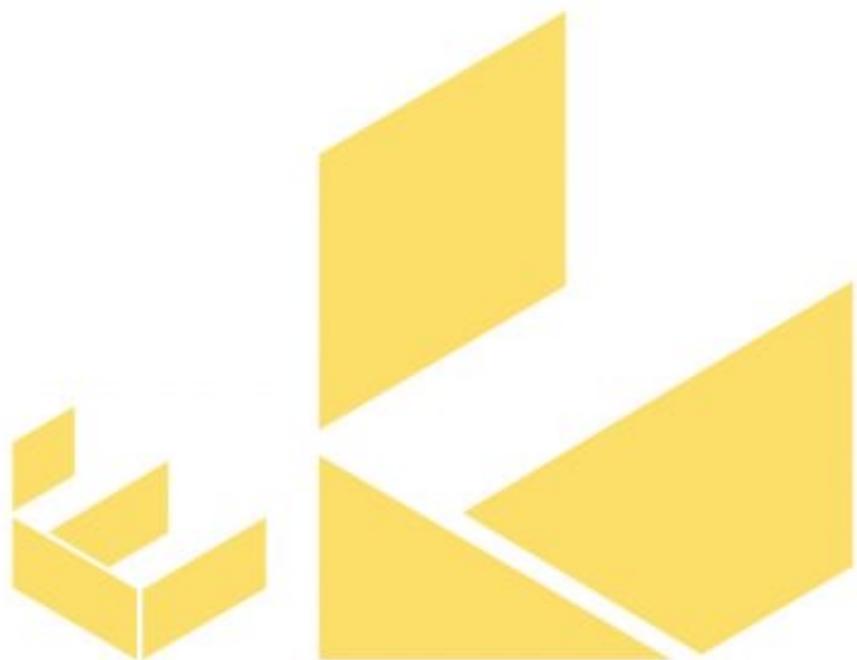


IMAGE / RASTER



Raster Graphics	Vector Graphics
They are composed of pixels.	They are composed of paths.
In Raster Graphics, refresh process is independent of the complexity of the image.	Vector displays flicker when the number of primitives in the image become too large.
Graphic primitives are specified in terms of end points and must be scan converted into corresponding pixels.	Scan conversion is not required.
Raster graphics can draw mathematical curves, polygons and boundaries of curved primitives only by pixel approximation.	Vector graphics draw continuous and smooth lines.
Raster graphics cost less.	Vector graphics cost more as compared to raster graphics.
They occupy more space which depends on image quality.	They occupy less space.
File extensions: .BMP, .TIF, .GIF, .JPG	File Extensions: .SVG, .EPS, .PDF, .AI, .DXF

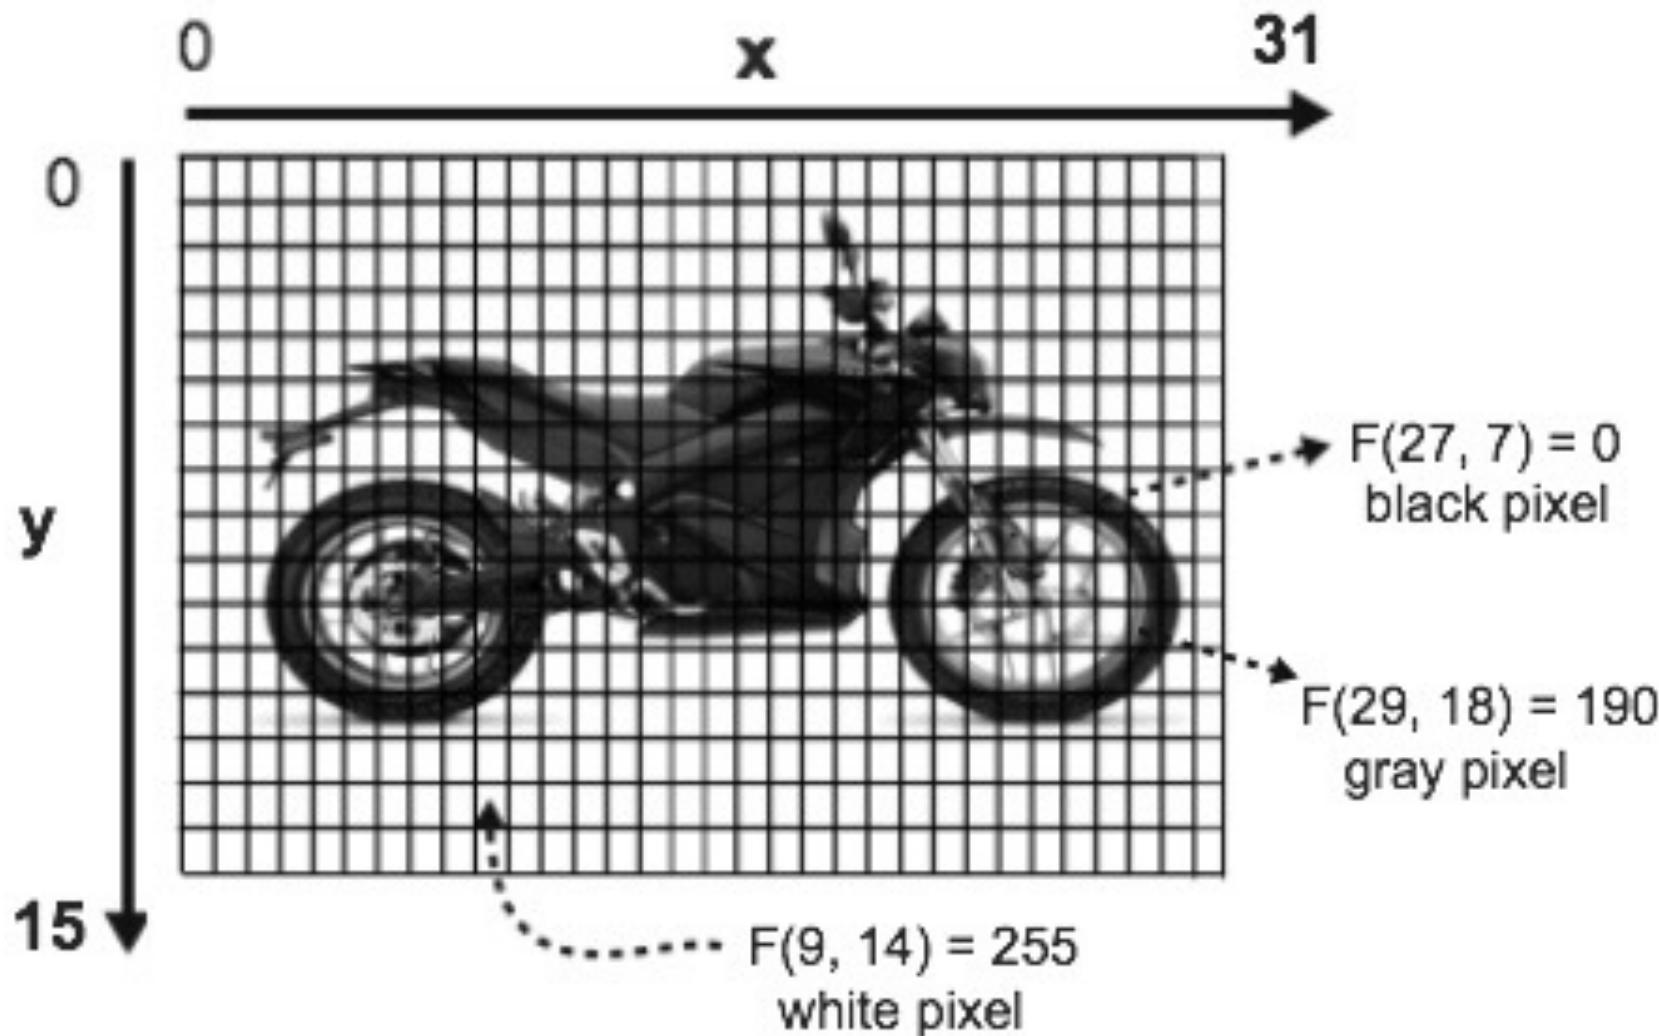
Image

An image is defined as a two-dimensional **function**, $F(x,y)$, where x and y are spatial coordinates, and the amplitude of F at any pair of coordinates (x,y) is called the **intensity** of that image at that point. When x,y , and amplitude values of F are finite, we call it a **digital image**.

Image
= function

= array

Grayscale image (32 x 16)



$$f(x, y) = \begin{bmatrix} f(0, 0) & f(0, 1) & \dots & f(0, N-1) \\ f(1, 0) & f(1, 1) & \dots & f(1, N-1) \\ \vdots & \vdots & \ddots & \vdots \\ f(M-1, 0) & f(N-1, 1) & \dots & f(M-1, N-1) \end{bmatrix}$$



What we see

08 02 22 97 38 15 00 40 00 75 04 05 07 78 52 12 50 77 91 08
49 49 99 40 17 81 18 57 40 87 17 40 98 43 69 48 04 54 62 00
81 49 31 73 55 79 14 29 93 71 40 67 53 88 30 03 49 13 36 65
52 70 95 23 04 60 11 42 69 24 68 54 01 32 56 71 37 02 36 91
22 31 16 71 51 67 43 89 41 92 36 54 22 60 40 28 66 33 13 80
24 47 32 60 99 03 45 02 44 75 33 53 78 36 84 20 35 17 12 50
32 98 81 28 44 23 47 10 26 38 40 67 59 54 70 66 18 38 64 70
67 26 20 68 02 62 12 20 95 63 98 39 63 08 40 91 66 69 94 21
24 55 58 05 66 73 99 26 97 17 78 78 96 83 16 88 36 89 63 72
21 36 23 09 75 00 76 46 20 45 35 14 00 61 33 97 36 31 33 95
78 17 53 29 22 75 31 67 15 94 03 80 04 62 16 14 09 53 56 92
16 59 05 42 94 35 31 47 55 58 88 24 00 17 54 24 36 29 85 57
86 56 00 48 35 71 89 07 05 44 44 37 44 40 21 58 51 54 17 58
19 80 81 68 05 94 47 69 28 73 92 13 86 52 17 77 04 89 55 40
04 52 08 83 97 35 99 16 07 97 57 32 16 26 26 79 33 27 98 66
88 36 48 87 57 62 20 72 03 46 33 67 46 55 12 32 63 93 53 69
04 42 14 73 38 25 39 11 24 94 72 18 08 46 29 32 40 62 76 36
20 69 36 41 72 30 23 88 34 62 99 69 82 47 59 85 74 04 36 16
20 73 35 29 78 31 90 01 74 31 49 71 48 86 81 14 23 57 05 54
01 70 54 71 83 51 54 49 16 92 33 48 61 43 52 01 89 19 67 48

What computers see



RGB channels

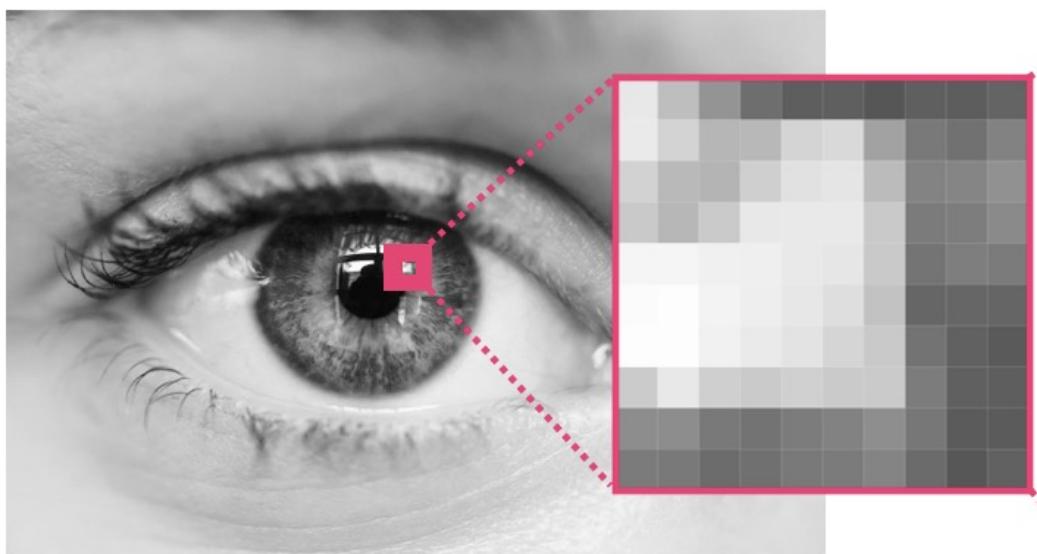
Channel 3 Blue intensity values					
35	165	163	165	158	...
166	166	164	166	159	...
156	159	162	165	159	...
160	169	167	169	169	...
170	170	168	170	170	...
160	162	166	169	170	...
11	158	156	158	158	...
159	159	157	159	159	...
149	151	155	158	159	...
146	146	149	153	158	...
145	143	143	148	158	...
...

Channel 2 Green intensity values					
102	169	167	169	169	...
170	170	168	170	170	...
160	162	166	169	170	...
170	170	168	170	170	...
160	162	166	169	170	...
11	158	156	158	158	...
159	159	157	159	159	...
149	151	155	158	159	...
146	146	149	153	158	...
145	143	143	148	158	...
...

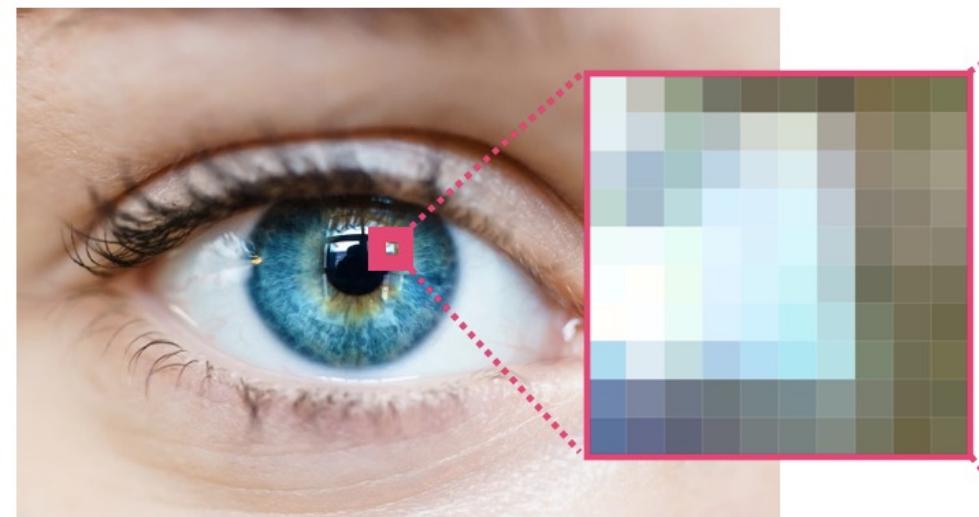
Channel 1 Red intensity values					
11	158	156	158	158	...
159	159	157	159	159	...
149	151	155	158	159	...
146	146	149	153	158	...
145	143	143	148	158	...
...

$$\text{Red} \quad + \quad \text{Green} \quad + \quad \text{Blue} \quad = \quad \text{Forest Green} \quad (11, 102, 35)$$

11 102 35



230	194	147	108	90	98	84	96	91	101
237	206	188	195	207	213	163	123	116	128
210	183	180	205	224	234	188	122	134	147
198	189	201	227	229	232	200	125	127	135
249	241	237	244	232	226	202	116	125	126
251	254	241	239	230	217	196	102	103	99
243	255	240	231	227	214	203	116	95	91
204	231	208	200	207	201	200	121	95	95
144	140	120	115	125	127	143	118	92	91
121	121	108	109	122	121	134	106	86	97



233	188	137	96	90	95	63	73	73	82
237	202	159	120	105	110	88	107	112	121
226	191	147	110	101	112	98	123	110	119
221	191	176	182	203	214	169	144	133	145
185	160	161	184	205	223	186	137	147	161
181	174	189	207	206	215	194	136	142	151
246	237	237	231	208	206	192	122	143	144
254	254	241	224	199	192	181	99	122	117
239	248	232	207	187	182	184	110	114	110
193	215	193	167	158	164	181	114	112	111
113	119	110	111	113	123	135	120	108	106
93	97	91	103	107	111	122	112	104	114

An image can be defined by a two-dimensional array specifically arranged in rows and columns.

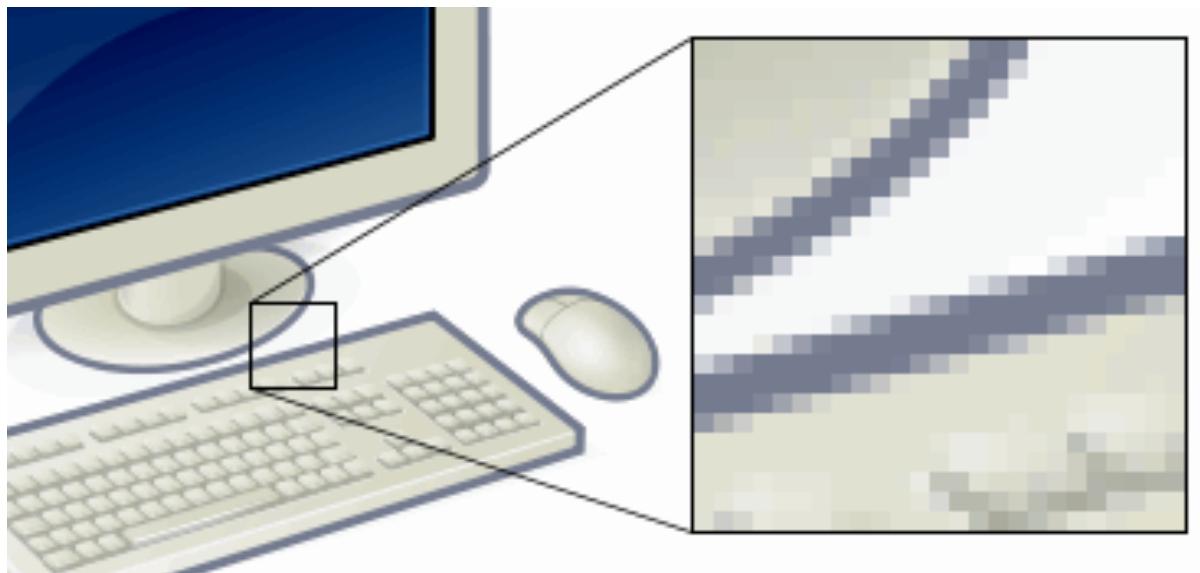
Digital Image is composed of a finite number of elements, each of which elements have a particular value at a particular location. These elements are referred to as *picture elements, image elements, and pixels*. A *Pixel* is most widely used to denote the elements of a Digital Image.



Copyright: Wikimedia commons,
GFDL and cc-by-sa-2.5,2.0,1.0

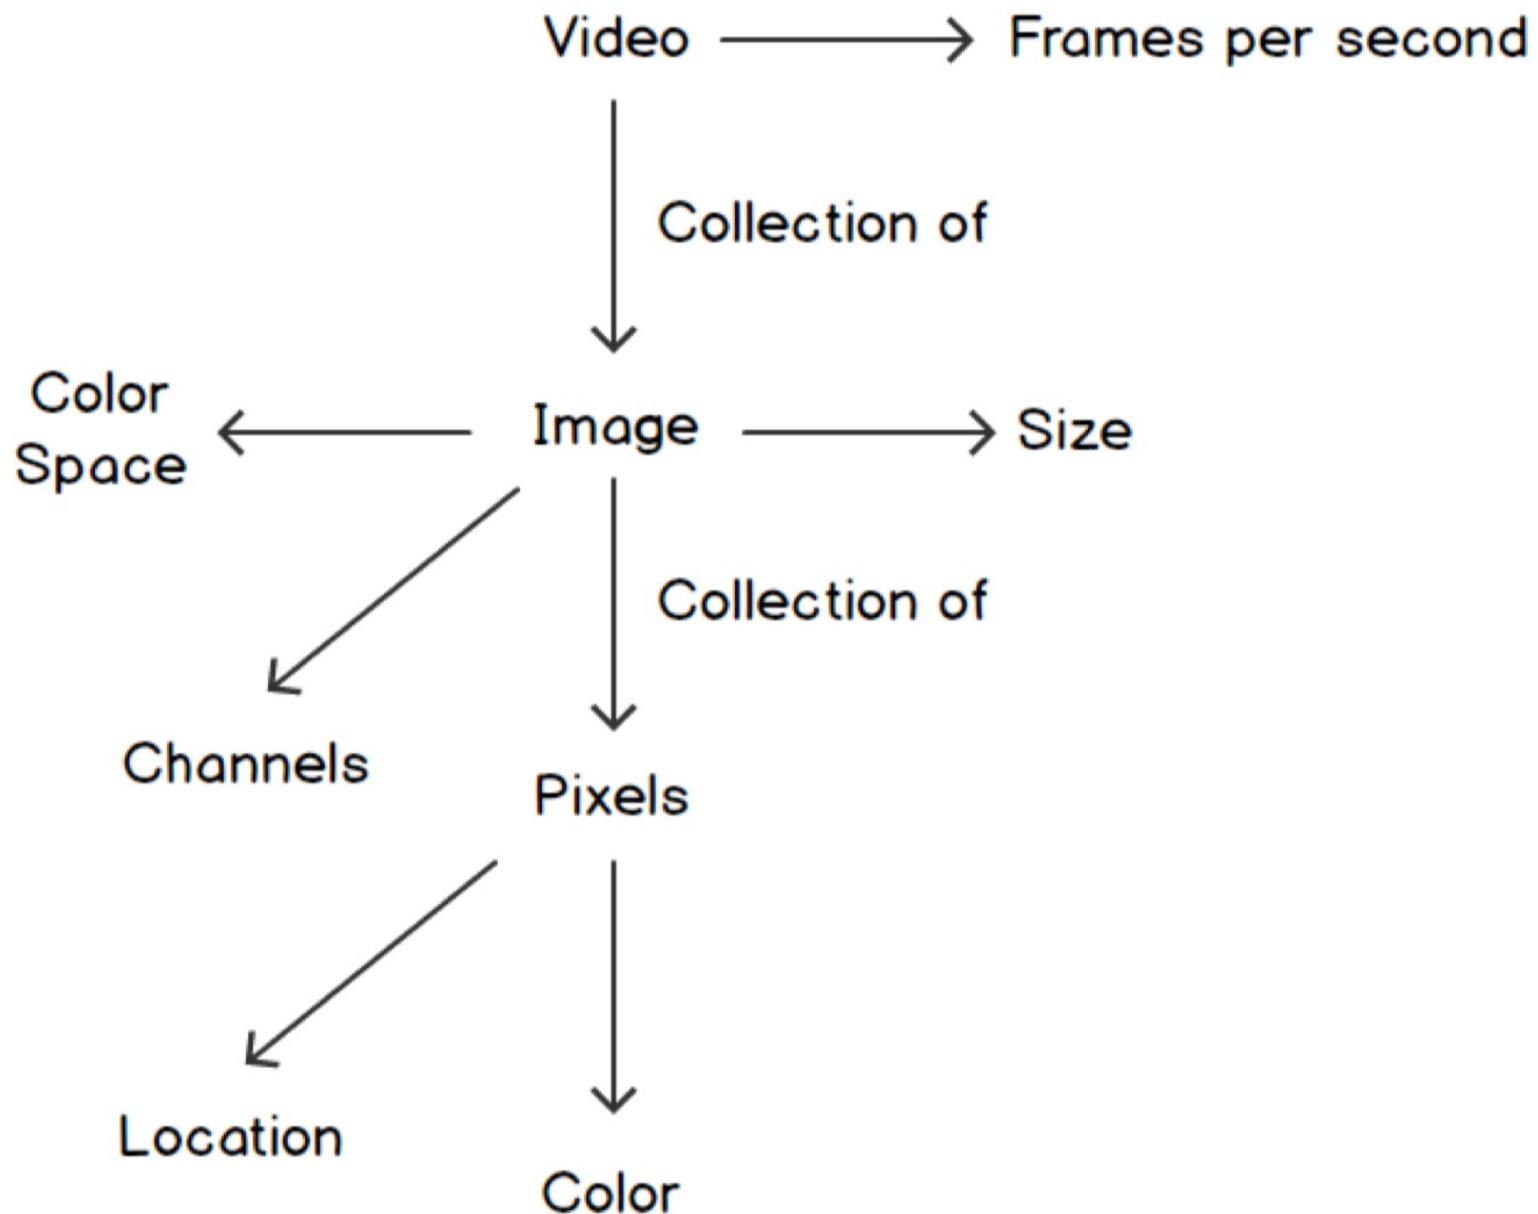
- PPI

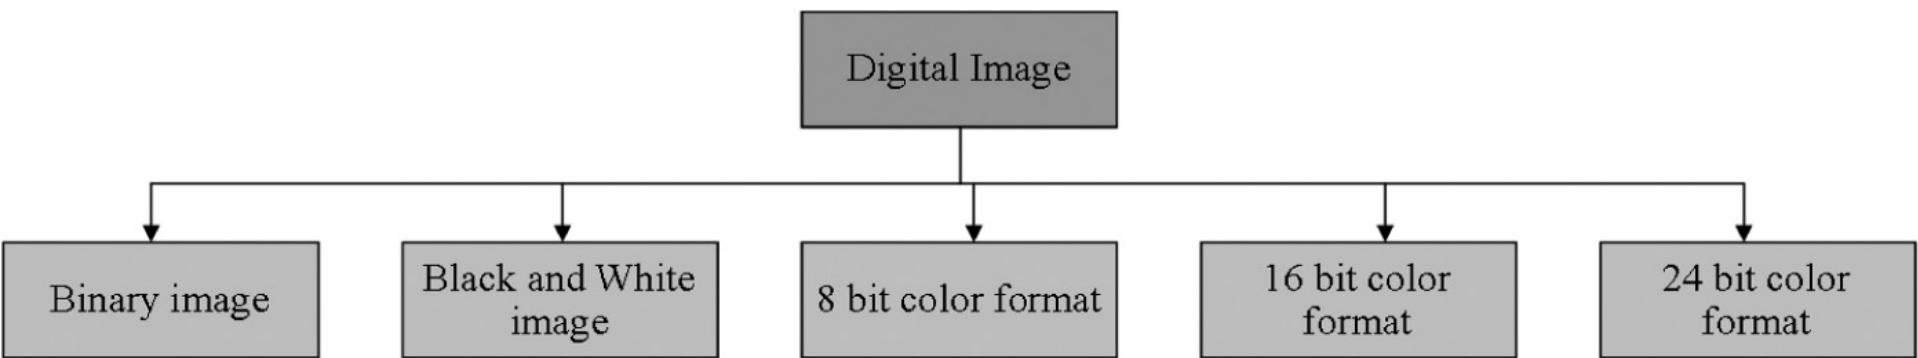
Pixels in Images



Copyright: Wikimedia commons,
GFDL and cc-by-sa-2.5,2.0,1.0

- 픽셀은 단색의 정사각형
- 전체 이미지의 크기를 표현할 때는 (세로픽셀수 x 가로픽셀수) 형식으로 표현
- The term pixel is generally applied to 2D and is replaced by voxel in 3D images.





■ Gray Scale

- 모든 색이 흑백. 각 픽셀은 명도를 나타내는 숫자로 표현. 0은 검은색을 나타내고 숫자가 커질수록 명도가 증가하여 하얀색이 됨. 숫자는 보통 0~255의 8비트 부호없는 정수로 저장

■ RGB

- Red, Green, Blue의 3가지 색의 명도를 뜻하는 숫자 3개가 합쳐진 벡터로 표현. 8비트 부호없는 정수를 사용하는 경우 (255, 0, 0)은 빨간색, (0, 255, 0)은 녹색, (0, 0, 255)는 파란색.
- 픽셀 데이터가 스칼라가 아닌 벡터이므로 이미지 데이터는 (세로픽셀수 x 가로픽셀수) 형태의 2차원 배열로 표현하지 못하고 (세로픽셀수 x 가로픽셀수 x 색채널) 형태의 3차원 배열로 저장

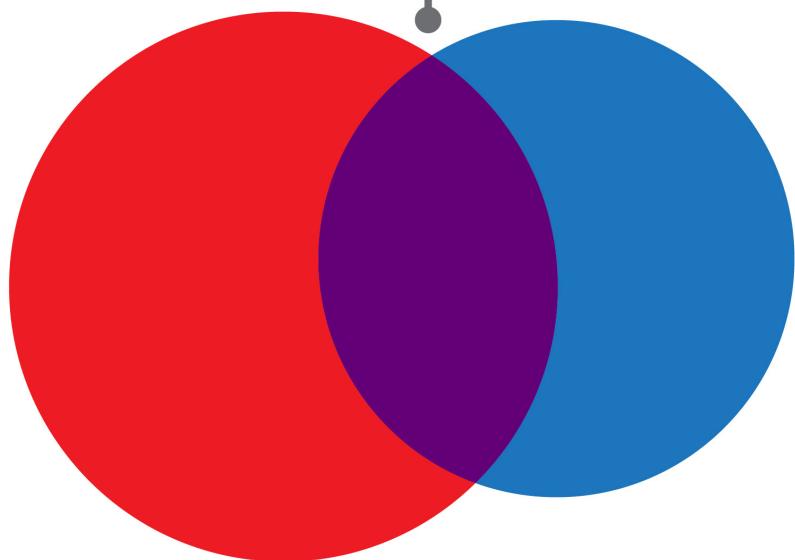
■ HSV

- Hue, Saturation, Value 세가지 값으로 표현
 - 색상(Hue): 색상값 H는 가시광선 스펙트럼을 주파수 별로 고리모양으로 배치했을 때의 각도. $0^\circ \sim 360^\circ$ 의 범위를 갖고 360° 와 0° 는 빨강을 가리킴
 - 채도(Saturation): 채도값 S는 특정한 색상의 진함의 정도. 가장 진한 상태를 100%이고 0%는 같은 명도의 무채색
 - 명도(Value): 명도값 V는 밝은 정도를 나타냄. 순수한 흰색, 빨간색은 100%이고 검은색은 0%
- HSV 색공간으로 표현된 파일은 imshow 명령으로 바로 볼 수 없음
 - Matplotlib : `rgb_to_hsv`, `hsv_to_rgb` 상호변환
- 색공간에 투명도(transparency)를 표현하는 A(Alpha) 채널이 추가된 RGBA, HSVA 등의 색 공간도 있음

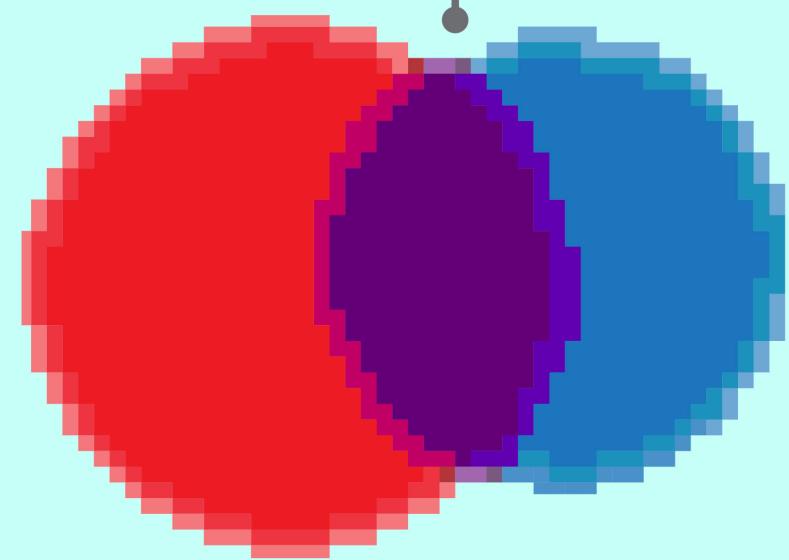
- **BMP** - Bitmap. This was probably the first type of digital image format that I can remember. Every picture on a computer seemed those days to be a BMP. In Windows XP the Paint program saves its images automatically in BMP. However, in Windows Vista and later images are now saved to JPEG. BMP is the basis platform for many other file types.
- **JPG / JPEG** - (Joint Photographic Experts Group) Jpeg format is used for color photographs, or any pictures with many blends or gradients. It is not good with sharp edges and tends to blur them a bit unless stored at high quality. This format became popular with the invention of the digital camera. Most, if not all, digital cameras download photos to your computer as a Jpeg file. Obviously the digital camera manufacturers see the value in high quality images that ultimately take up less space.
- **GIF** - (Graphics Interchange Format) Gif format is best used for text, line drawings, screenshots, cartoons, and animations. Gif is limited to a total number of 256 colors or less, so Gif images are relatively small. It is commonly used for fast loading web pages. It also makes a great banner or logo for your web-page. Animated pictures can also be saved in GIF format as a sequence of static images. For example, a flashing banner would be saved as a Gif file.
- **PNG** - (Portable Networks Graphic) This lossless formats is one of the best image formats. It was not always compatible with all web browsers or image software, but nowadays it is the best image format to use for website. I use .png for logos and screenshots. One of its most astonishing abilities is being able to compress images losslessly (without loss of pixels), although the final compressed size varies between image editors.
- **TIFF** - (Tagged Image File Format) This file format has not been updated since 1992 and is now owned by Adobe. It can store an image and data (tags) in the one file. TIFF can be compressed, but it is rather its ability to store image data in a lossless format that makes a TIFF file a useful image archive, because unlike standard JPEG files, a TIFF file using lossless compression (or none) may be edited and re-saved without losing image quality. This file is commonly used for scanning, faxing, word processing, and so on. It is no longer a common file format to use with your digital photos, as jpeg is great quality and takes up less space.

	BMP	JPG	PNG	GIF
Compressed	FALSE	TRUE	TRUE	TRUE
Lossless	TRUE	FALSE	TRUE	TRUE
Transparency	FALSE	FALSE	TRUE	TRUE
Translucency	FALSE	FALSE	TRUE	FALSE
Recommended for photographs	FALSE	TRUE	FALSE	FALSE
Recommended for static graphics/icons	FALSE	FALSE	TRUE	FALSE
Recommended for animated graphics/icons	FALSE	FALSE	FALSE	TRUE

Vector



Raster



AI

EPS

CGM

PDF

SVG

CDR

BMP

TIFF

PCX

GIF

PNG

JPEG

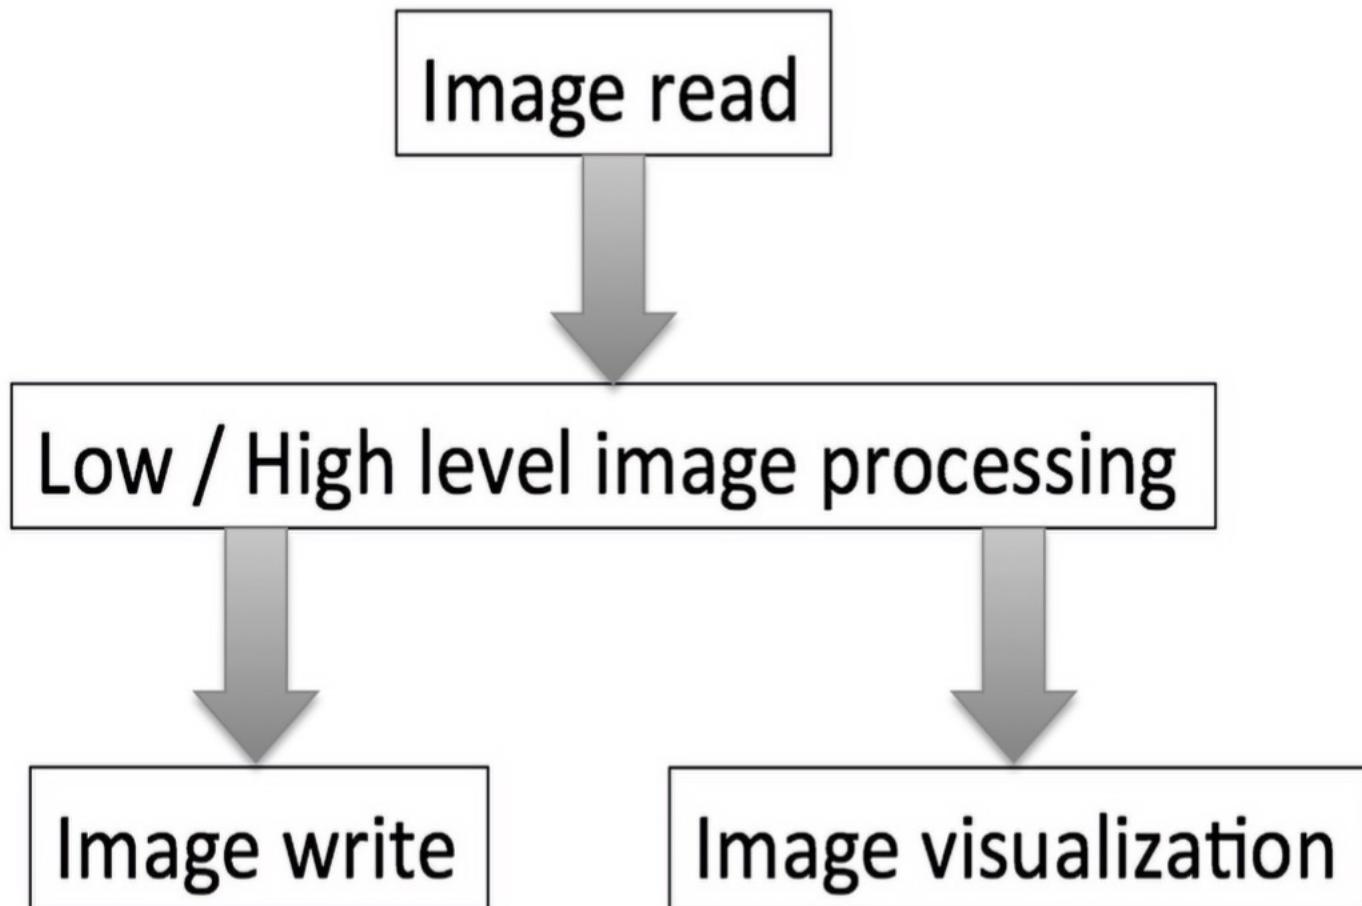
Image Processing

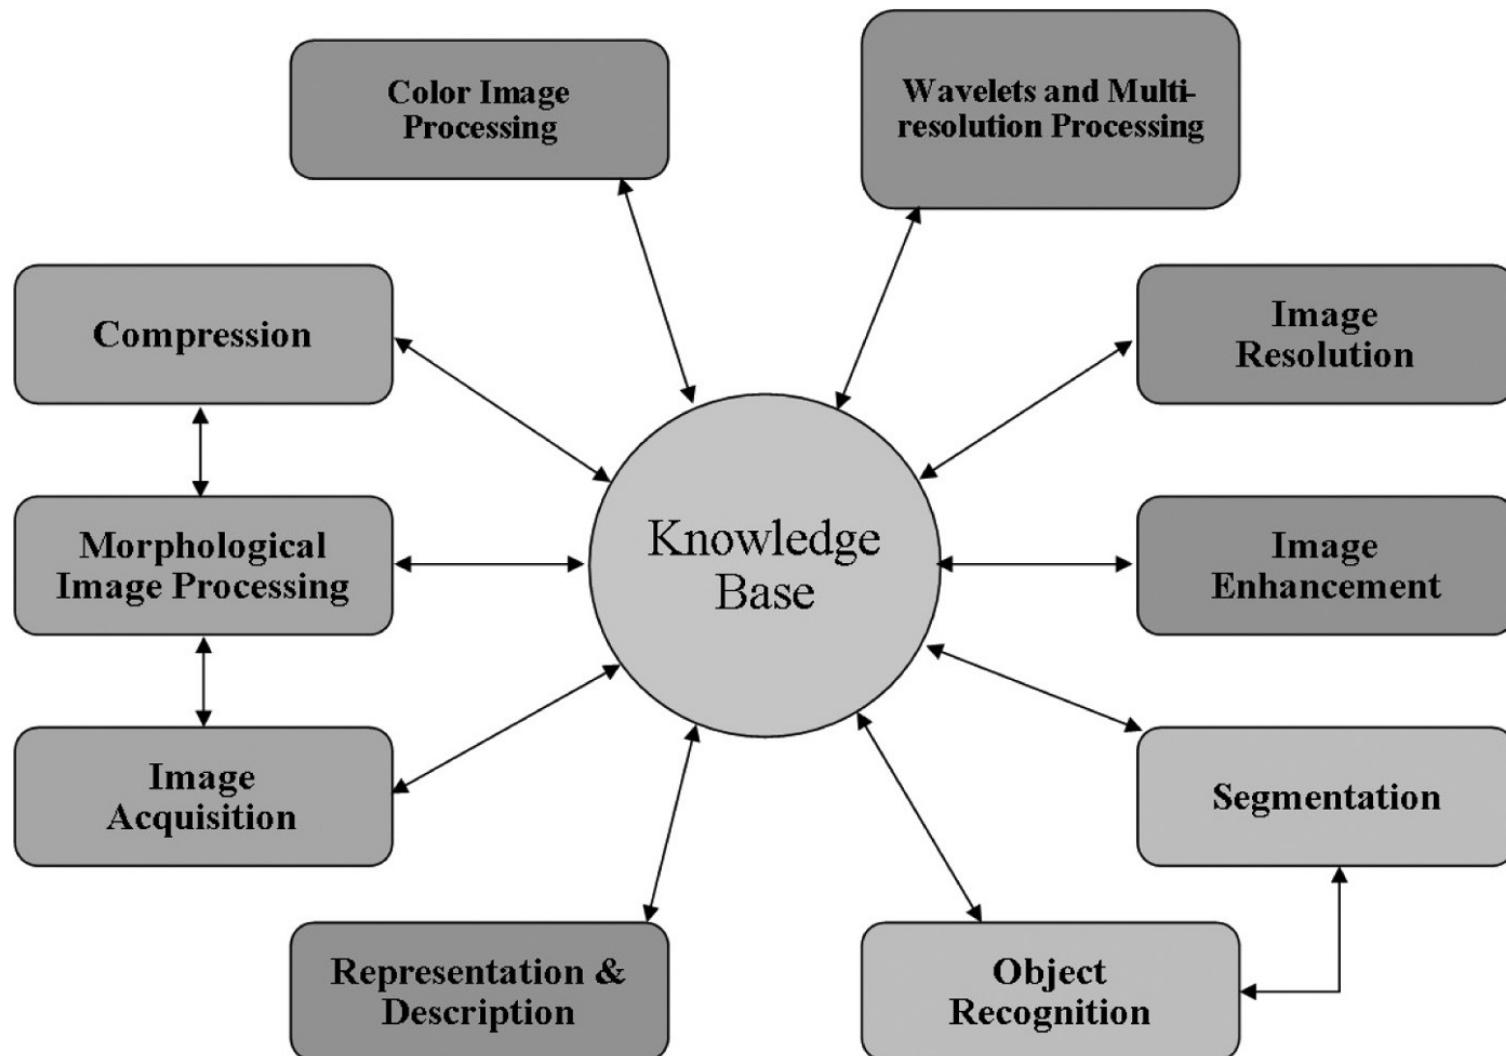
Digital Image Processing means processing digital image by means of a digital computer.

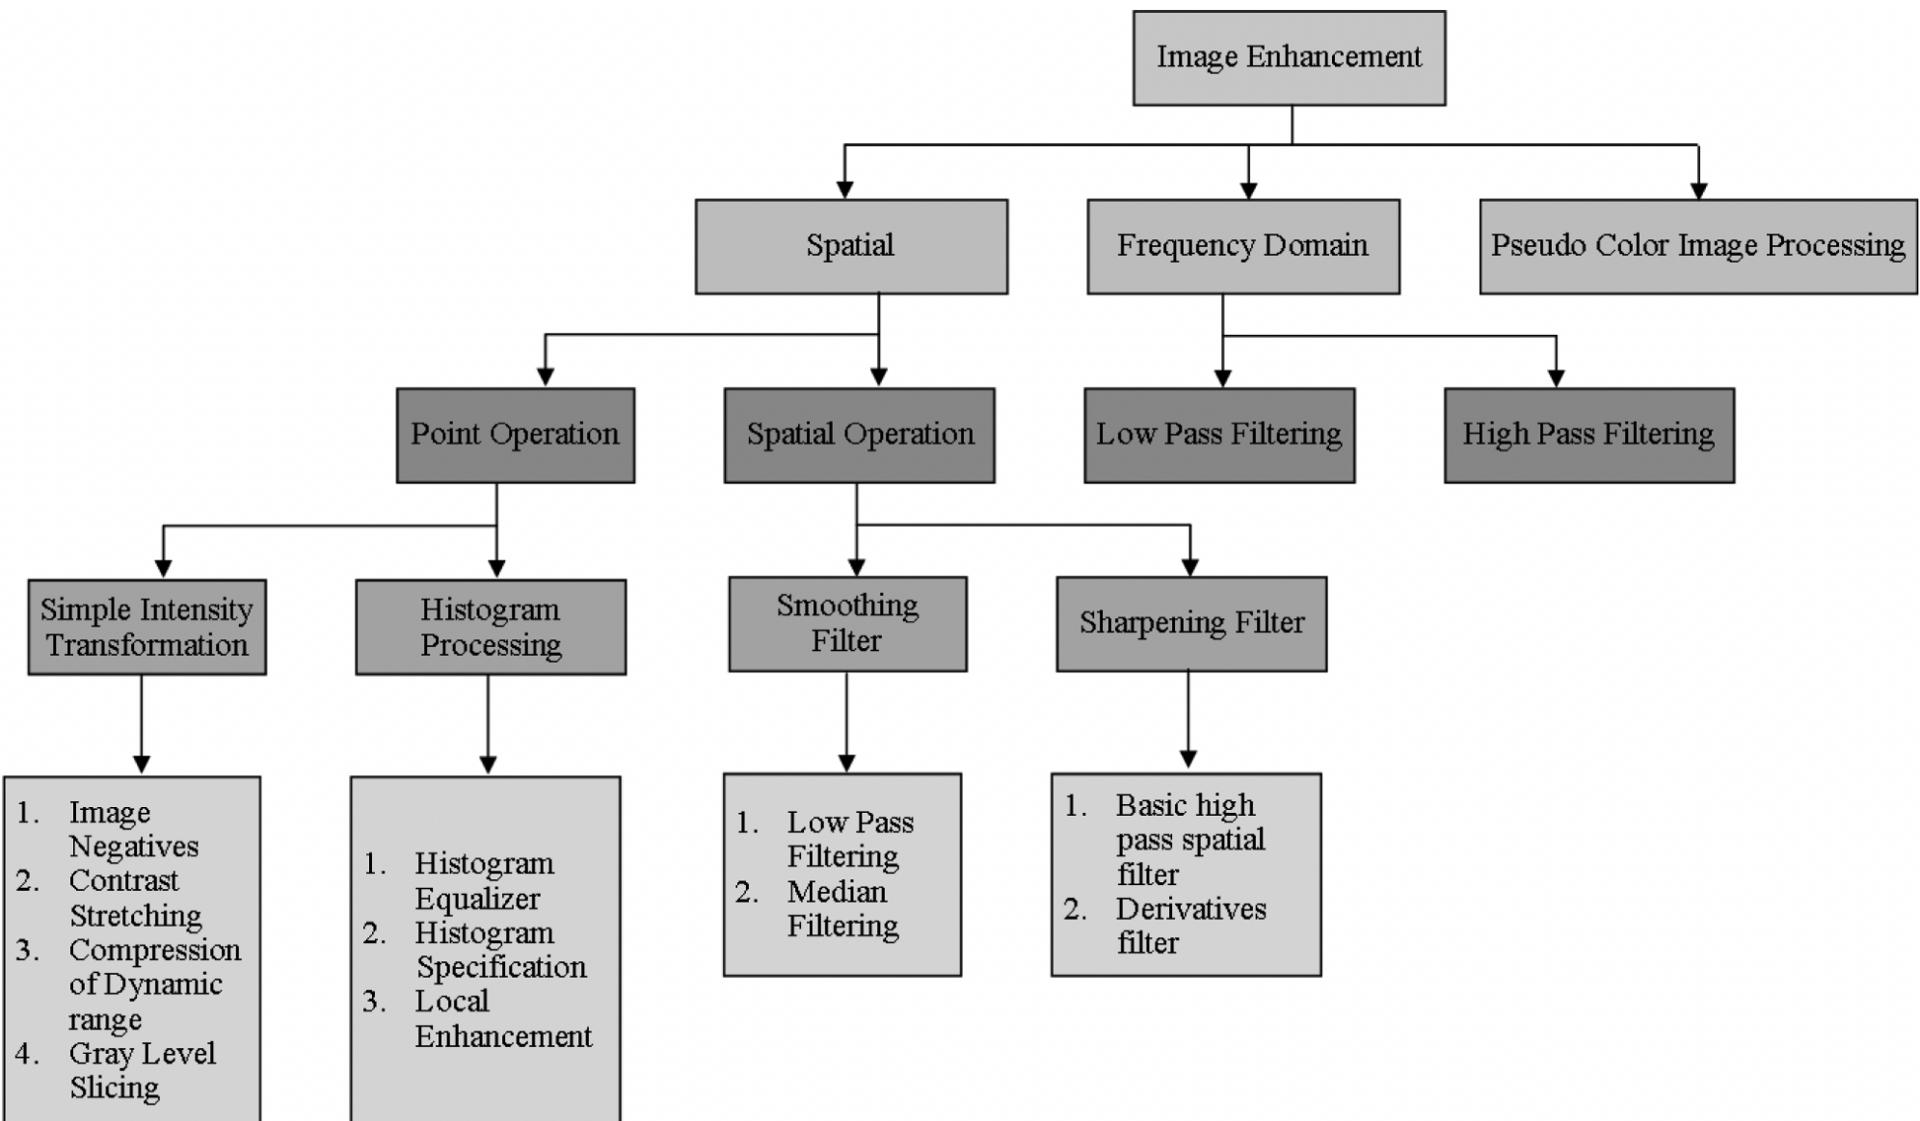
We can also say that **it is a use of computer algorithms**, in order to get enhanced image to extract some useful information.

Image processing mainly include the following steps:

1. Importing the image via image acquisition tools
2. Analysing and manipulating the image
3. Output in which result can be altered image or a report which is based on analysing that image.







No Free Lunch

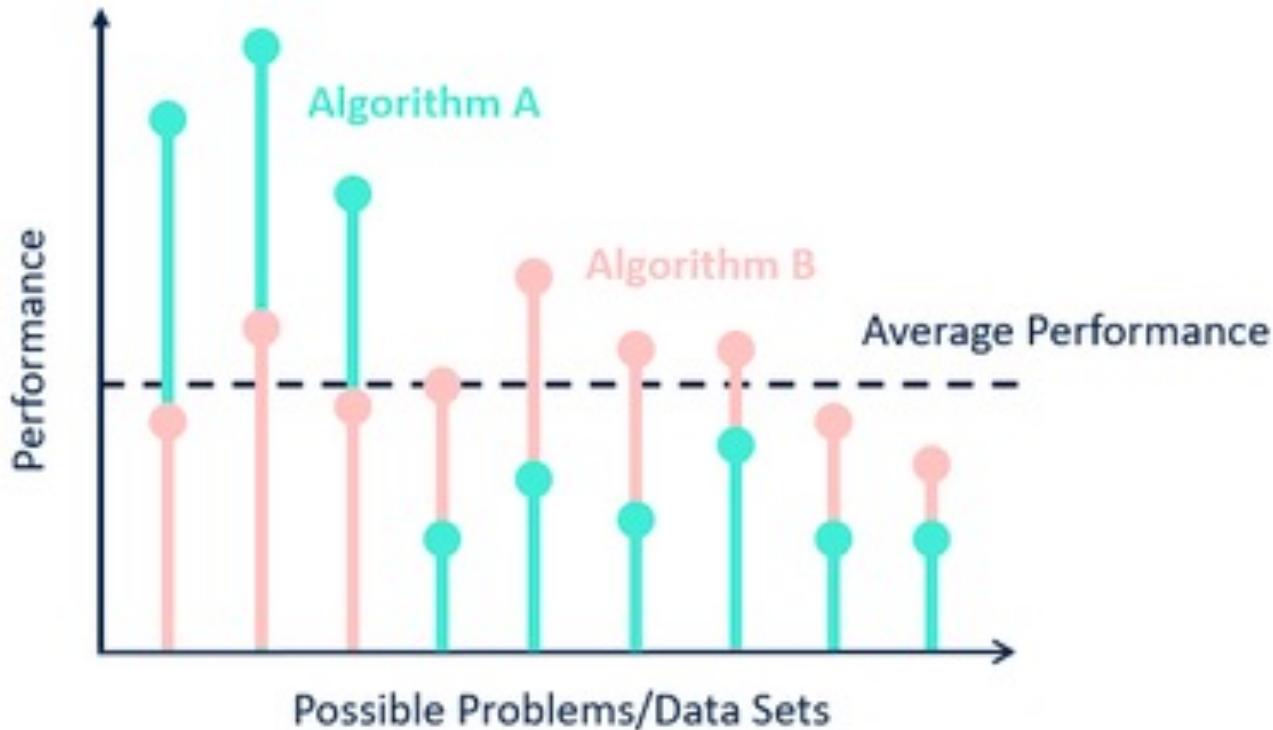
1. Our model is a simplification of reality
2. Simplification is based on assumptions (model bias)
3. Assumptions fail in certain situations

“No one model works best for all possible situations.”

The **supervised** learning no-free-lunch theorems. (2002)

D. H. Wolpert.

No Free Lunch Theorems for **Optimizations** (1997)



어떤 특정 정책에 의해 얼핏 보면 이득을 얻는 것 같지만,
그것은 한 측면의 이득일 뿐이고 반드시 이면에 다른 측면이 있고, 그 측면에서 손해가 발생

마스터 알고리즘은 우리의 미래를 어떻게 바꾸는가

페드로 도밍고스 저 / 김형진 옮김 / 최승진 감수

완벽한 마스터 알고리즘이 탄생하는 순간
세상 모든 것이 재편될 것이다!

인공지능과 무인자동차, HCI, 클라우드 컴퓨팅, 사용자 인터넷까지
우리 삶을 변화시킬 가장 혁신적인 기술, 머신러닝의 모든 것

데이터과학 분야의 최고 영예인 SIGKDD 혁신상을
2년 연속 수상한 세계 최고의 머신러닝 전문가
페드로 도밍고스가 전하는 머신러닝의 현재와 미래

아마존
컴퓨터·기술
1위

A Few Useful Things to Know about Machine Learning

Pedro Domingos
Department of Computer Science and Engineering
University of Washington
Seattle, WA 98195-2350, U.S.A.
pedrod@cs.washington.edu

ABSTRACT

Machine learning algorithms can figure out how to perform important tasks by generalizing from examples. This is often feasible and cost-effective where manual programming is not. As more data becomes available, more ambitious problems can be tackled. As a result, machine learning is widely used in computer science and other fields. However, developing successful machine learning applications requires a substantial amount of “black art” that is hard to find in textbooks. This article summarizes two key lessons that machine learning researchers and practitioners have learned. These include pitfalls to avoid, important issues to focus on, and answers to common questions.

1. INTRODUCTION

Machine learning systems automatically learn programs from data. This is often a very attractive alternative to manually constructing them, and in the last decade the use of machine learning has spread rapidly throughout computer science and beyond. Machine learning is used in Web search, spam filters, recommender systems, ad placement, credit scoring, fraud detection, stock trading, drug design, and many other applications. A recent report from the McKinsey Global Institute asserts that machine learning (a.k.a. data mining or predictive analytics) will be the driver of the next big wave of innovation. Several books are available for interested practitioners and researchers (e.g., [25]). However, much of the “folk knowledge” that is needed to successfully develop machine learning applications is not readily available in them. As a result, many machine learning projects take much longer than necessary or wind up producing less-than-ideal results. Yet much of this folk knowledge is fairly easy to communicate. This is the purpose of this article.

Many different types of machine learning exist, but for illustration purposes I will focus on the most mature and widely used one: classification. Nevertheless, the issues I will discuss apply across all of machine learning. A *classifier* is a system that inputs (typically) a vector of discrete and/or continuous *feature values* and outputs a single discrete value, the *class*. For example, it can filter spam messages into “spam” or “not spam,” and its input may be a Boolean vector $\mathbf{x} = (x_1, \dots, x_j, \dots, x_n)$, where $x_j = 1$ if the j th word in the dictionary appears in the email and $x_j = 0$ otherwise. A *learner* inputs a *training set of examples* (\mathbf{x}_i, y_i) , where $\mathbf{x}_i = (x_{i,1}, \dots, x_{i,n})$ is an observed input and y_i is the corresponding output, and outputs a classifier. The test of the learner is whether this classifier produces the

correct output y_i for future examples \mathbf{x}_t (e.g., whether the spam filter correctly classifies previously unseen emails as spam or not spam).

2. LEARNING = REPRESENTATION + EVALUATION + OPTIMIZATION

Suppose you have an application that you think machine learning might be good for. The first problem facing you is to decide what kind of machine learning is applicable. Which one to use? There are literally thousands available, and hundreds more are published each year. The key to not getting lost in this huge space is to realize that it consists of combinations of just three components. The components are:

Representation. A classifier must be represented in some formal language that the computer can handle. Conversely, choosing a representation for a task is tantamount to choosing the set of classifiers that it can possibly learn. This set is called the *hypothesis space* of the learner. If a classifier is not in the hypothesis space, it cannot be learned. A related question, which we will address in a later section, is how to represent the input, i.e., what features to use.

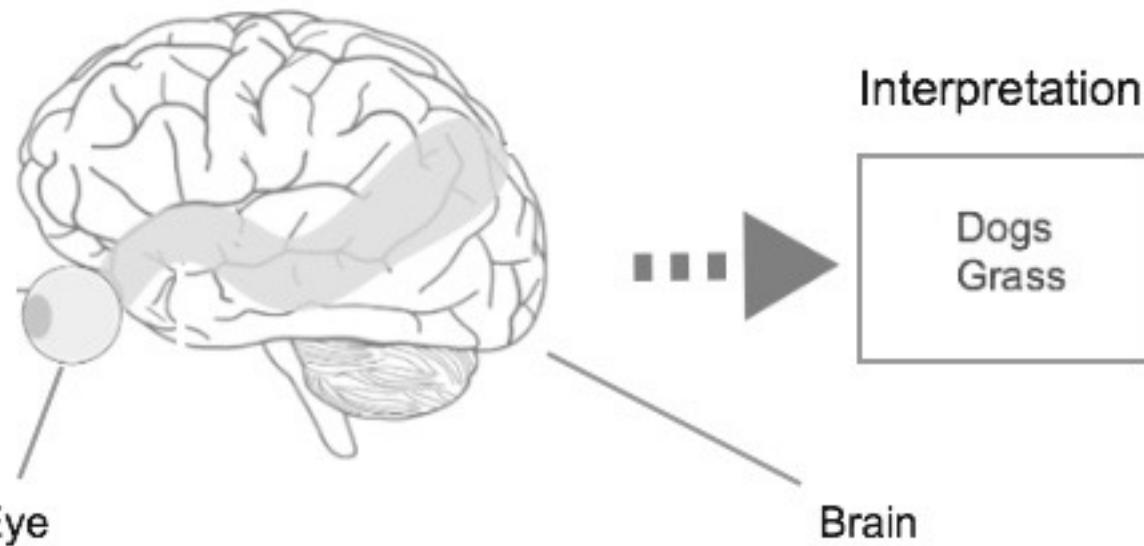
Evaluation. An evaluation function (also called *objective function* or *scoring function*) is needed to distinguish good classifiers from bad ones. The evaluation function used internally by the algorithm may differ from the external one that we want the classifier to optimize, for ease of optimization (see below) and due to the issues discussed in the next section.

Optimization. Finally, we need a method to search among the classifiers in the language for the highest-scoring one. The choice of optimization technique is key to the efficiency of the learner, and also helps determine the classifier produced if the evaluation function has more than one optimum. It is common for new learners to start out using off-the-shelf optimizers, which are later replaced by custom-designed ones.

Table 1 shows common examples of each of these three components. For example, *k-nearest neighbor* classifies a test example by finding the k most similar training examples and predicting the majority class among them. Hyperplane-based methods form a linear combination of the features per class and predict the class with the highest-valued combination. Decision trees test one feature at each internal node,

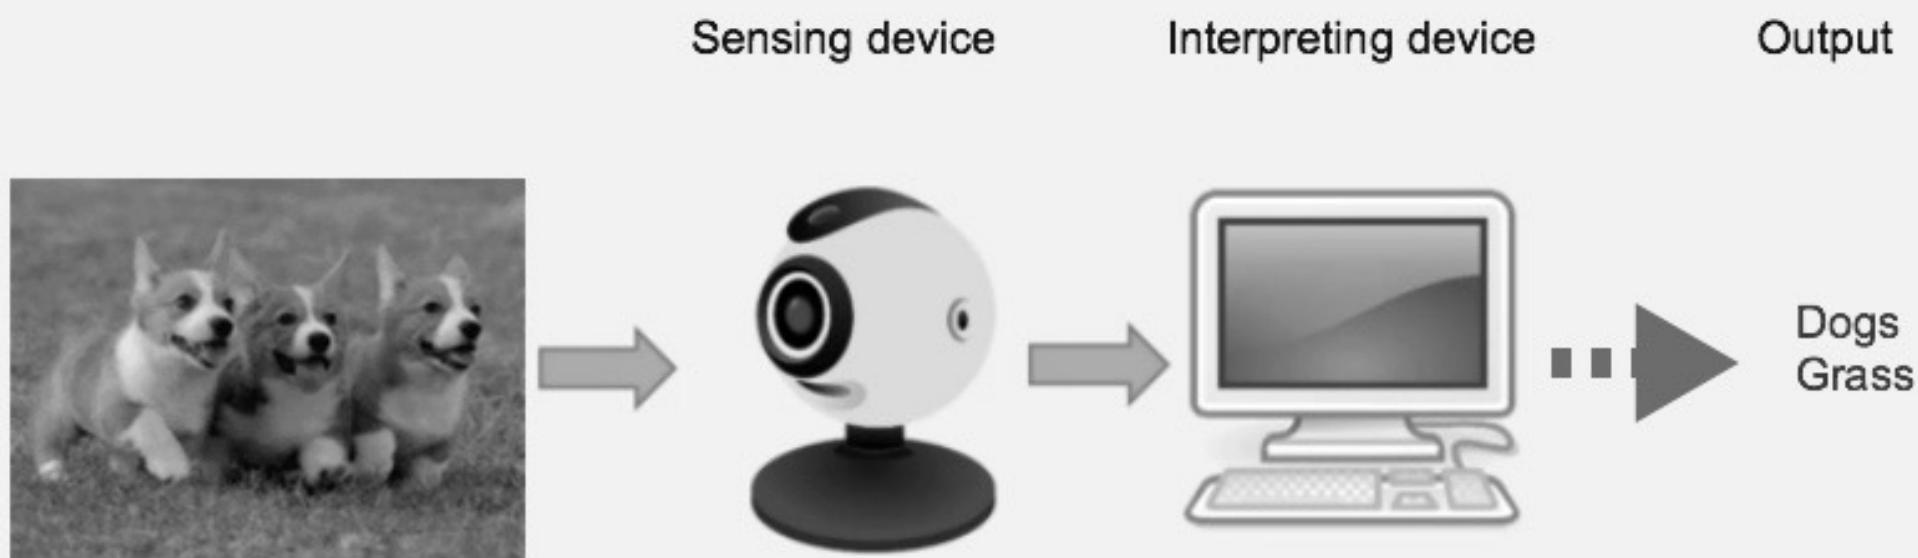


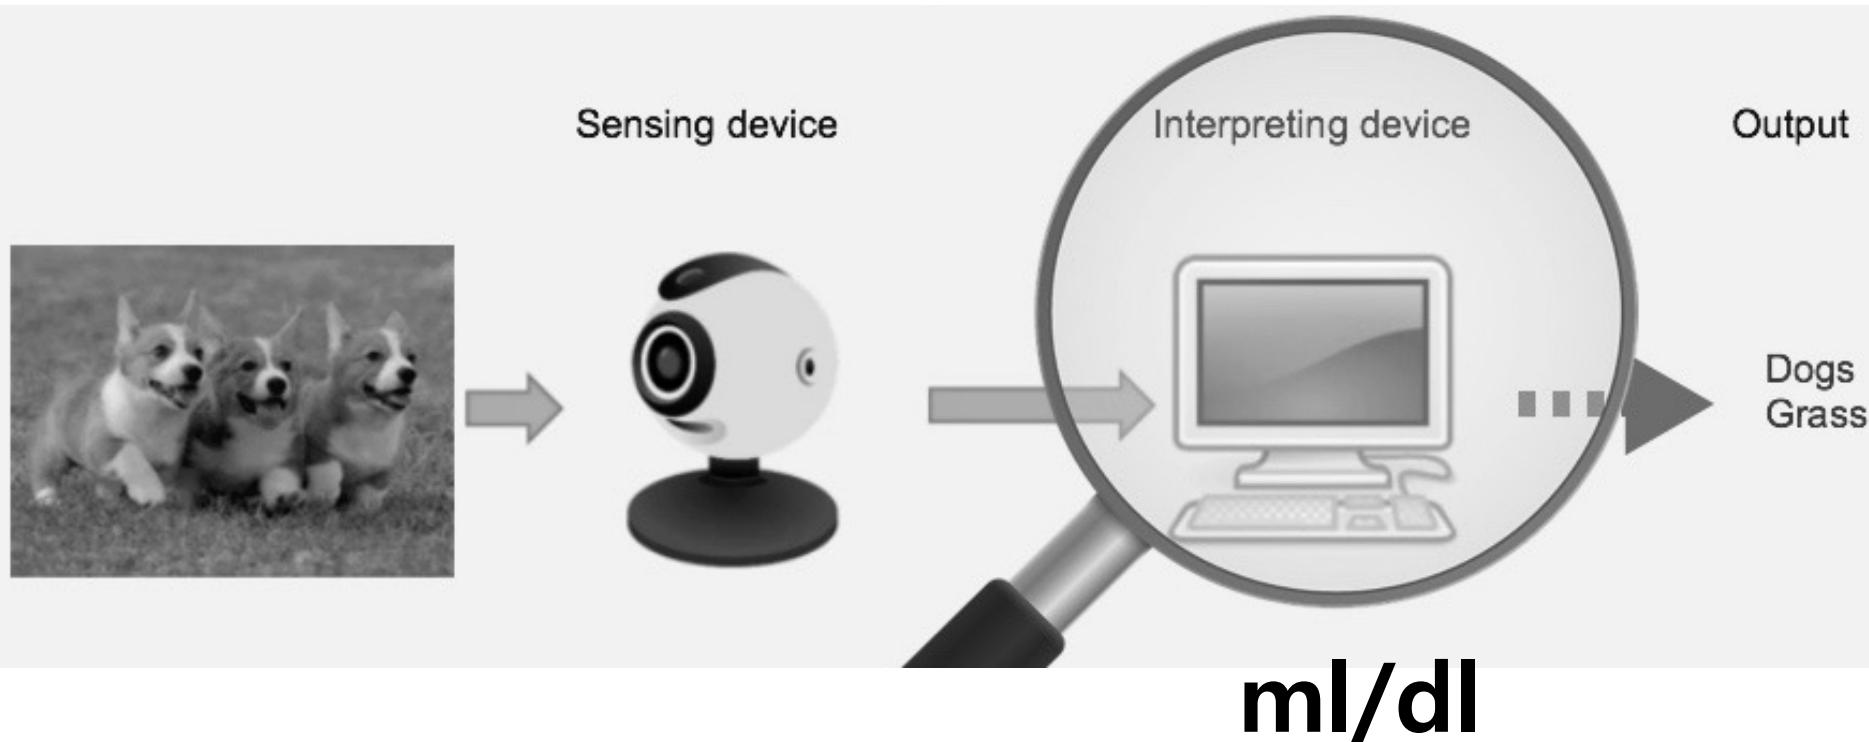
Human Vision System

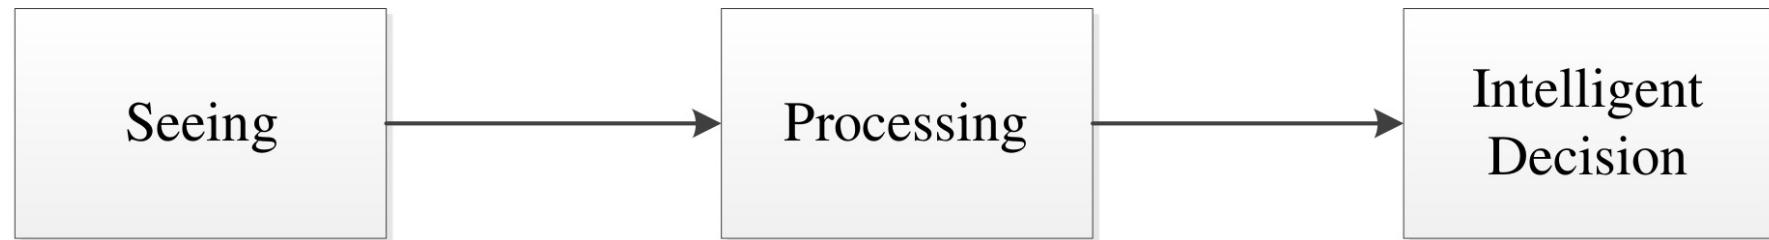


(sensing device responsible for
capturing images of the environment)

(Interpreting device responsible for
understanding the image content)



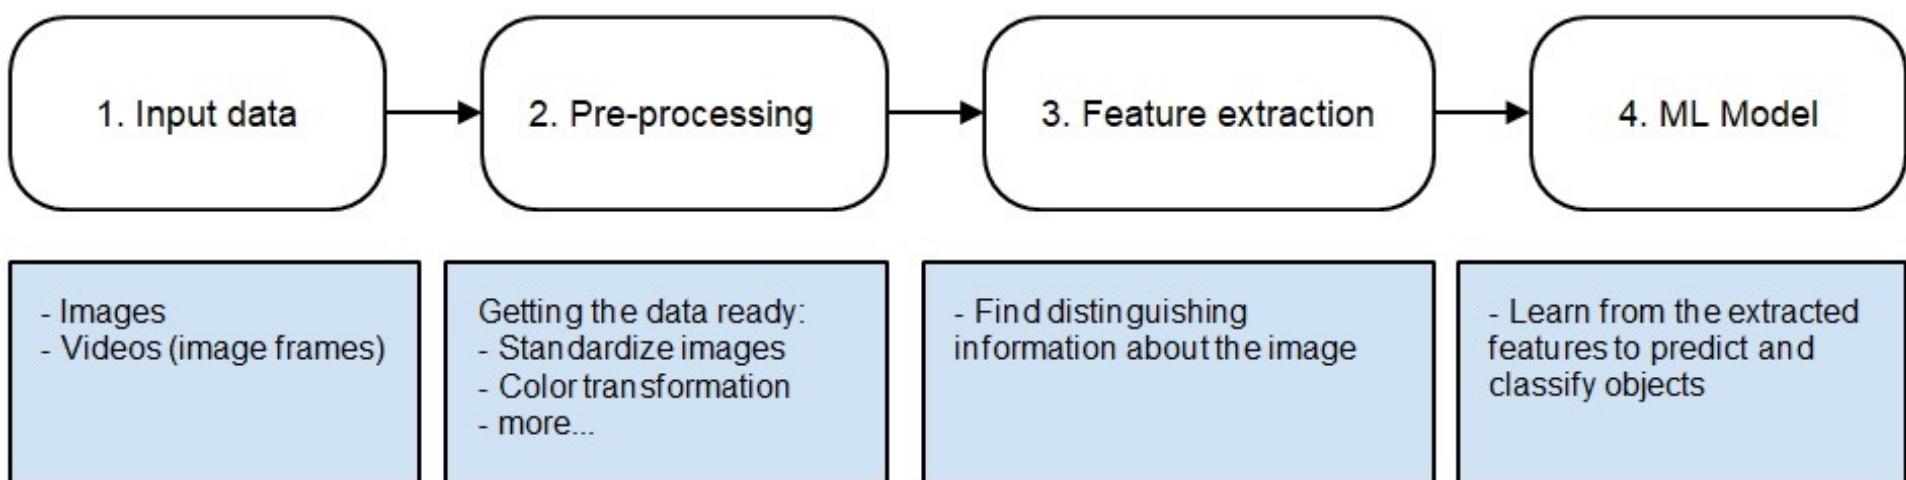


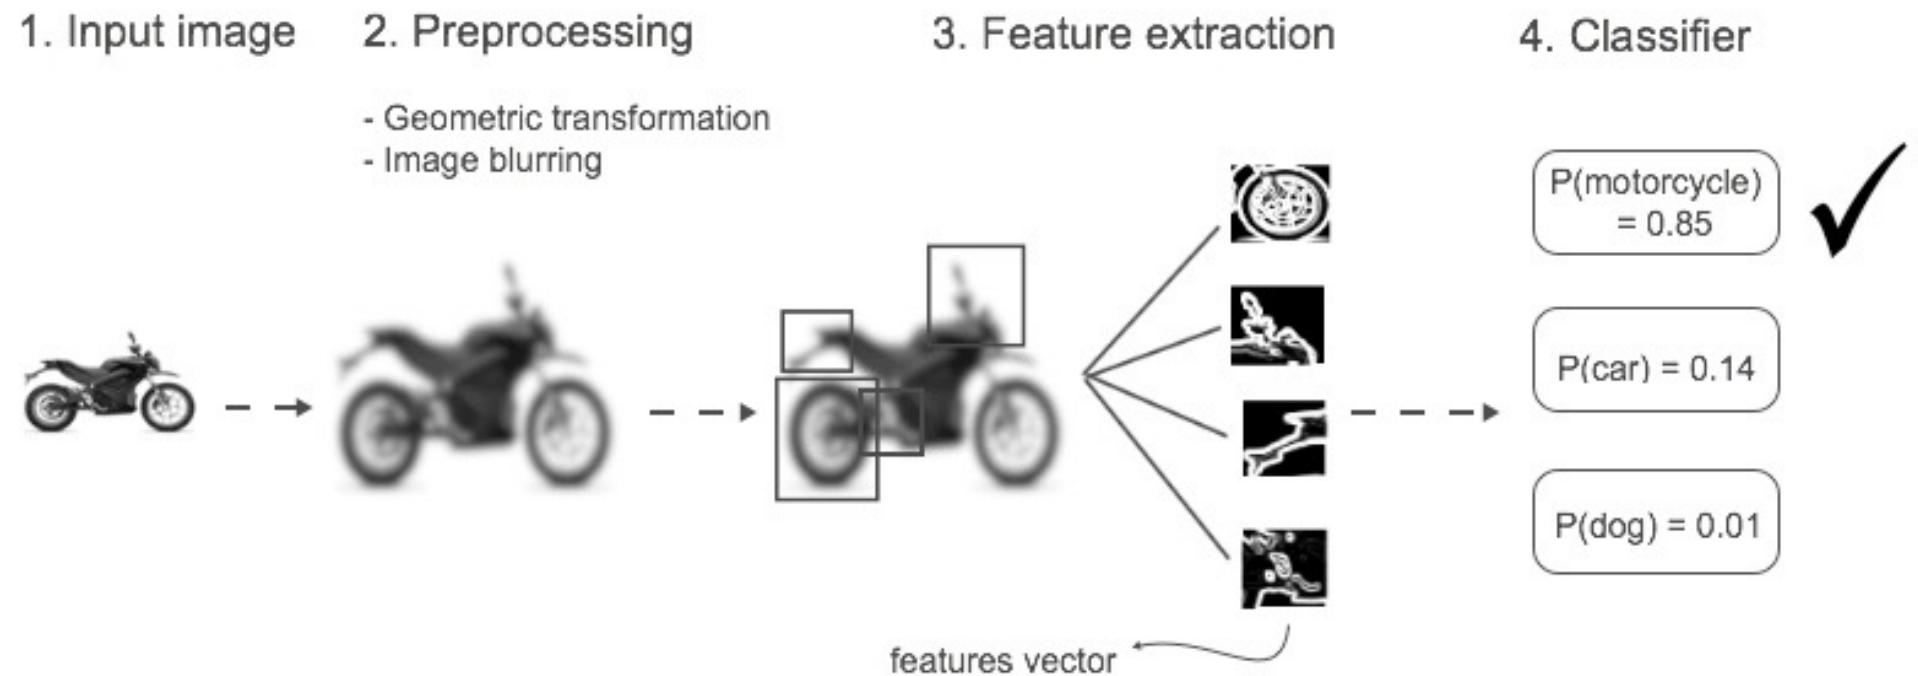


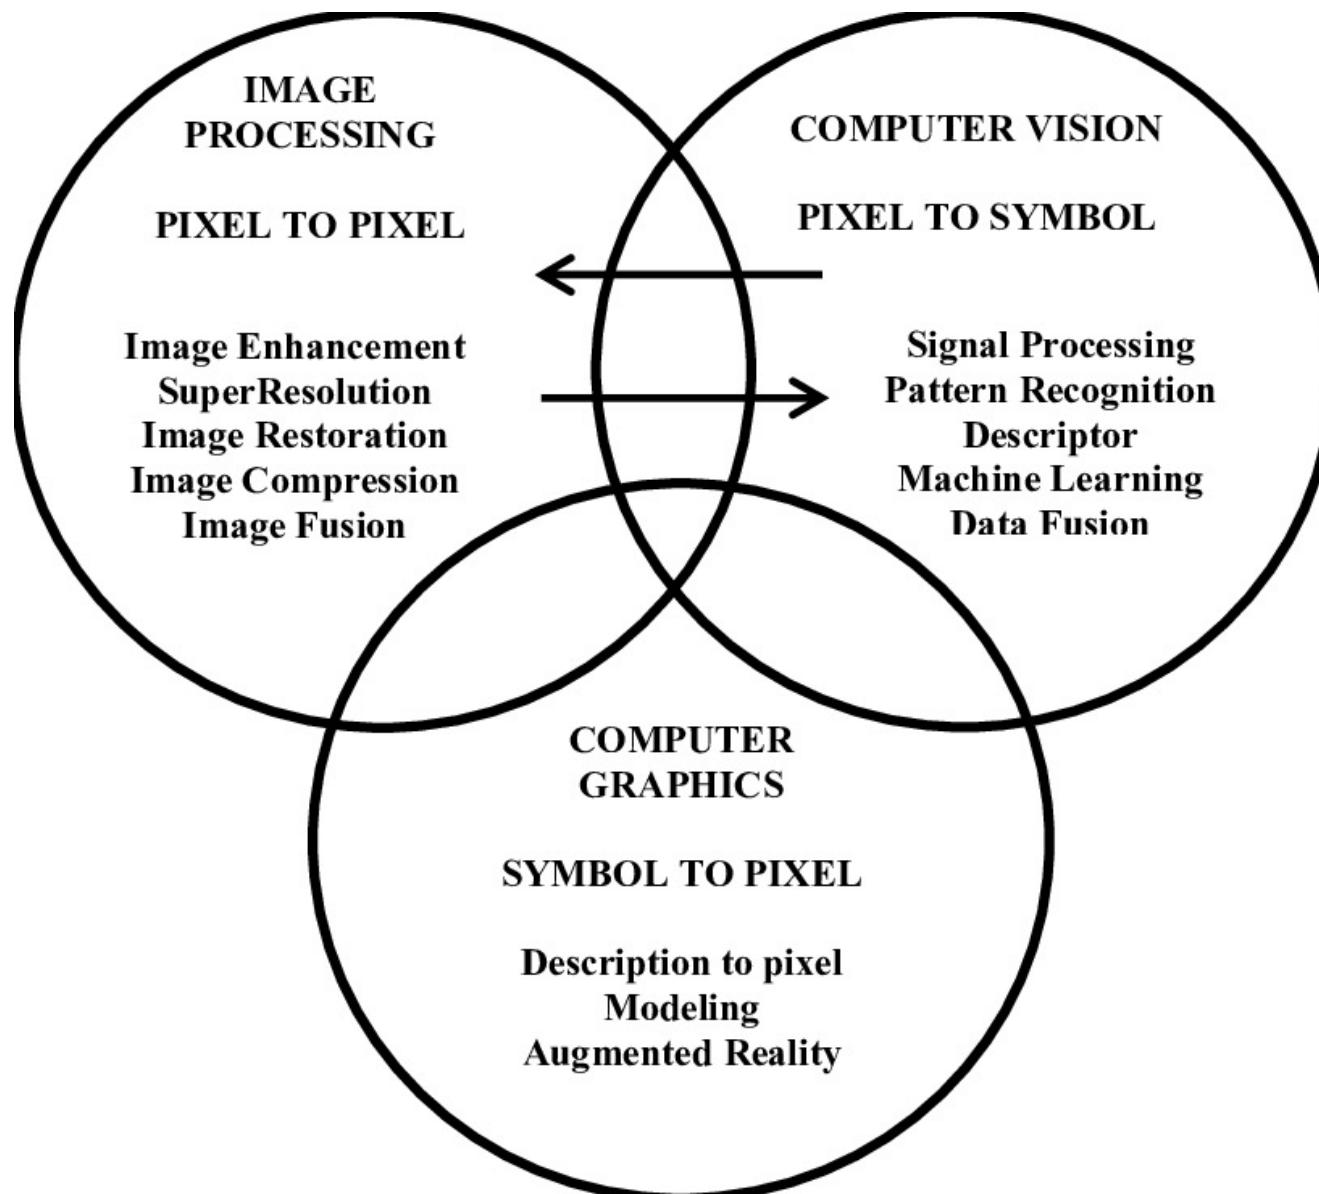
(a)



(b)







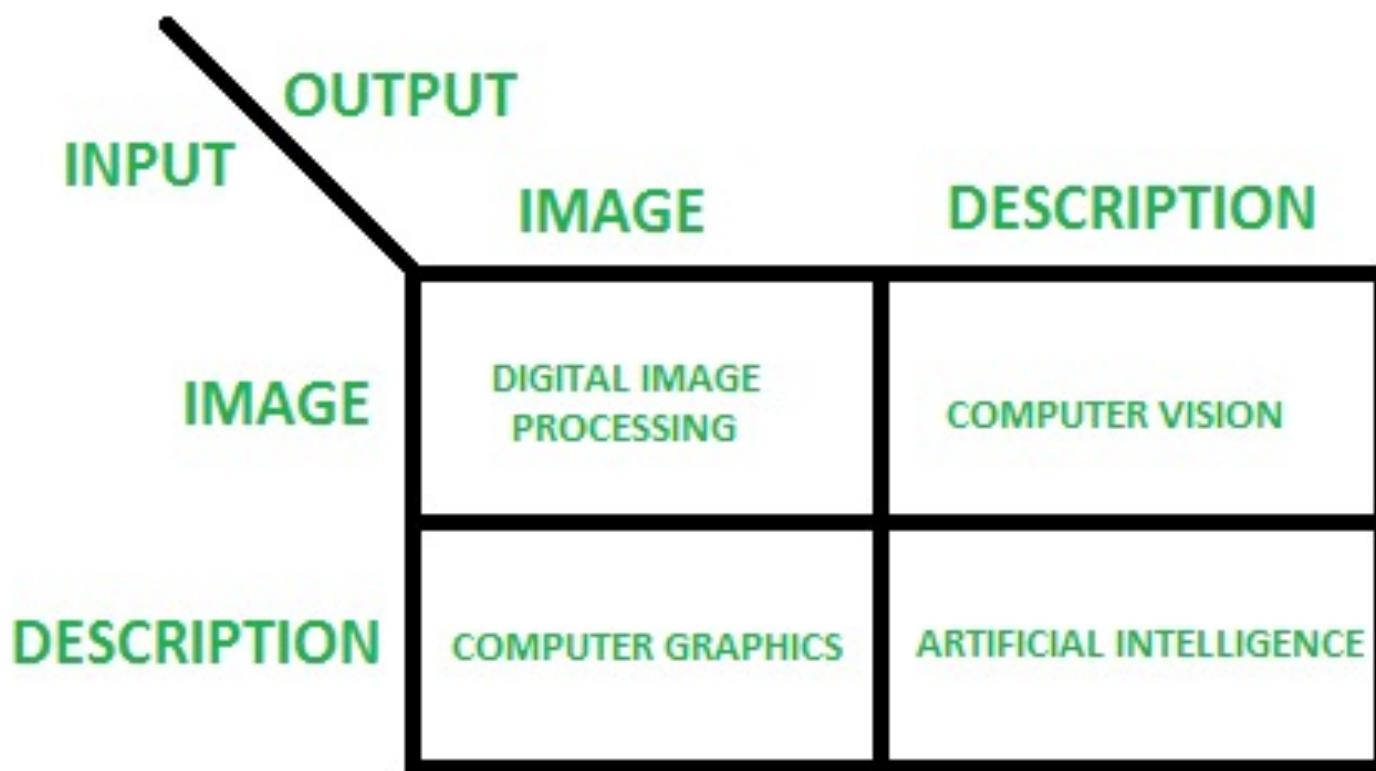


Image processing is the process of creating a new image from an existing image, typically simplifying or enhancing the content in some way. It is a type of digital signal processing and is not concerned with understanding the content of an image.

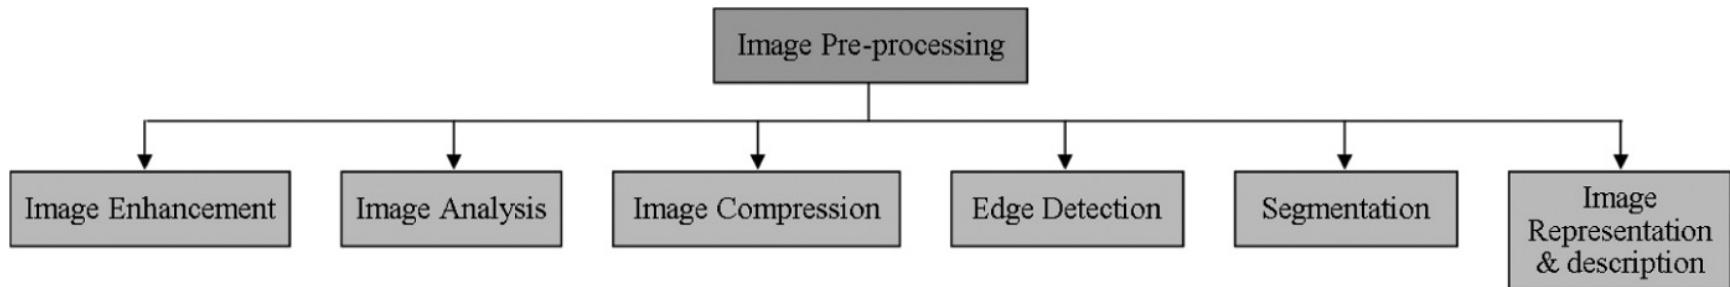
A given computer vision system may require image processing to be applied to raw input, e.g. pre-processing images.

Examples of image processing include:

- Normalizing photometric properties of the image, such as brightness or color.
- Cropping the bounds of the image, such as centering an object in a photograph.
- Removing digital noise from an image, such as digital artifacts from low light levels.

Image Processing for ML/DL

■ Transitional Image Pre-processing



■ Essential

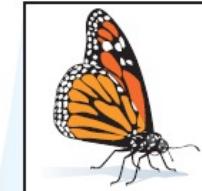
- Uniform Image Size

■ Techniques to improve performance of machine learning models

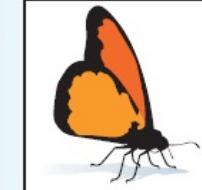
- Uniform Aspect Ratio
- Normalized Image Inputs
- Uniform Image Size
- Dimensionality Reduction
 - Convert color images to grayscale to reduce computation complexity
- Mean and Perturbed Images
 - Standardize images
- Data Augmentation



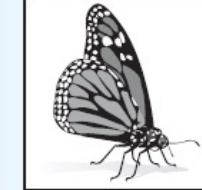
Data augmentation



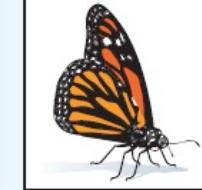
Original image



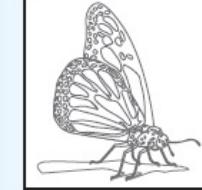
De-texturized



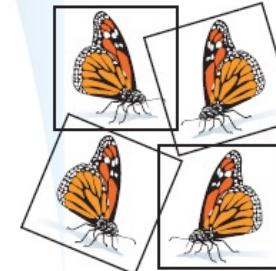
De-colorized



Edge enhanced



Salient edge map



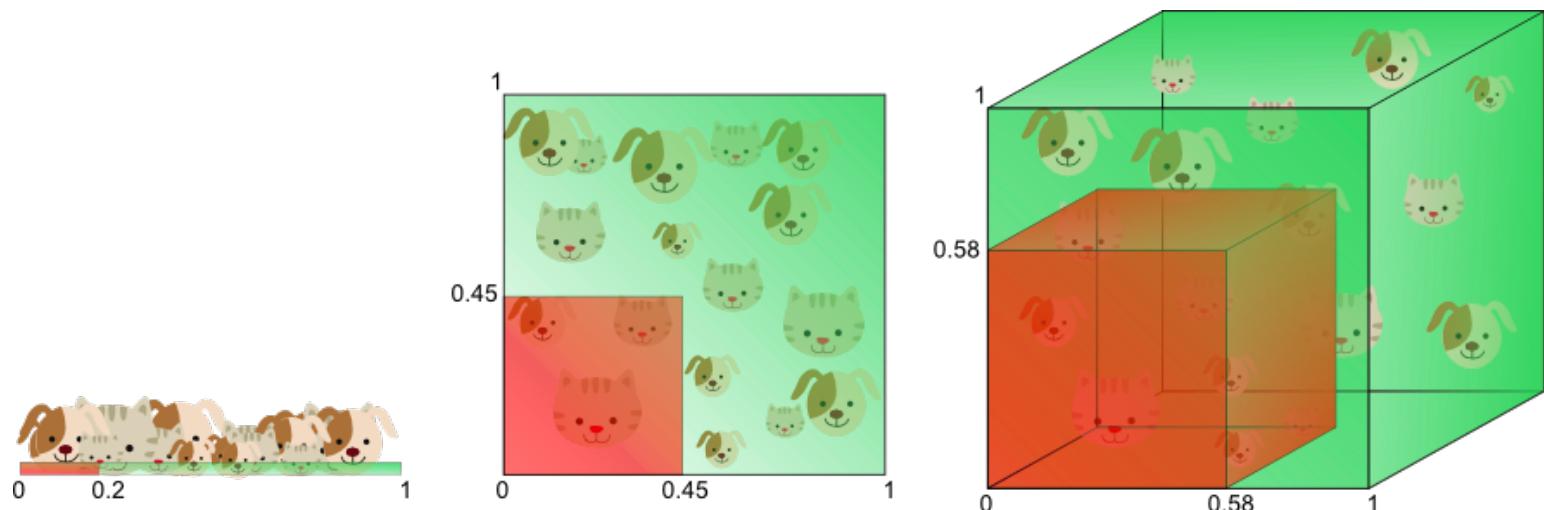
Flip/rotate

Why use features?

데이터의 차원이 증가할수록

학습에 필요한 데이터의 개수가 기하급수적으로 늘어나는 문제

- 같은 비율의 공간을 모델링하기 위해서 필요한 데이터 량이 급속도록 증가
- Example: 공간의 20%를 채우기 위해서 필요한 데이터 수
- $0.2(1\text{차원}) < 0.9(2\text{차원}) < 1.74 (3\text{차원})$



■ The Curse of Dimensionality

- 저차원에서는 충분했던 데이터 양도 고차원이 되면 공간(학습 모델)을 설명하기에 부족해 질 수 있음
- 모델을 설명하기에 데이터가 부족 > Overfitting 문제

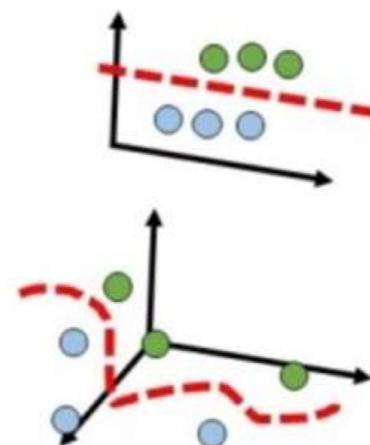
■ 해결책

- 데이터 양
- 중요한 특징만 뽑자
 - Dimensionality Reduction
 - Feature Extraction
 - Feature Selection
- 사전지식 이용 > Regularization



Good

Over Fitting
(차원의 저주)



충분성: 충분하지 않는 데이터

대표성: sampling noise : 우연에 의한 대표성 없는 데이터

sampling bias : 큰 샘플도 표본추출방법이 잘못되면 대표성을 띠지 못함

[데이터 품질]

불완전성: 데이터에 특성이 없거나 값이 누락

노이즈가 많은: 데이터에 잘못된 레코드(손상된) 또는 이상값이 있음

불일치: 데이터에 충돌하는 레코드 또는 일치하지 않는 값

Garbage in, Garbage out

- 우수한 예측 모델을 구축하려면 우수한 데이터가 필요
- 데이터 품질을 높여서 궁극적으로 모델 성능을 높이기 위해서 데이터 상태 검사를 수행하여 조기에 데이터 문제를 발견하고 적절한 데이터 처리 및 정리 단계를 결정하는 것이 중요

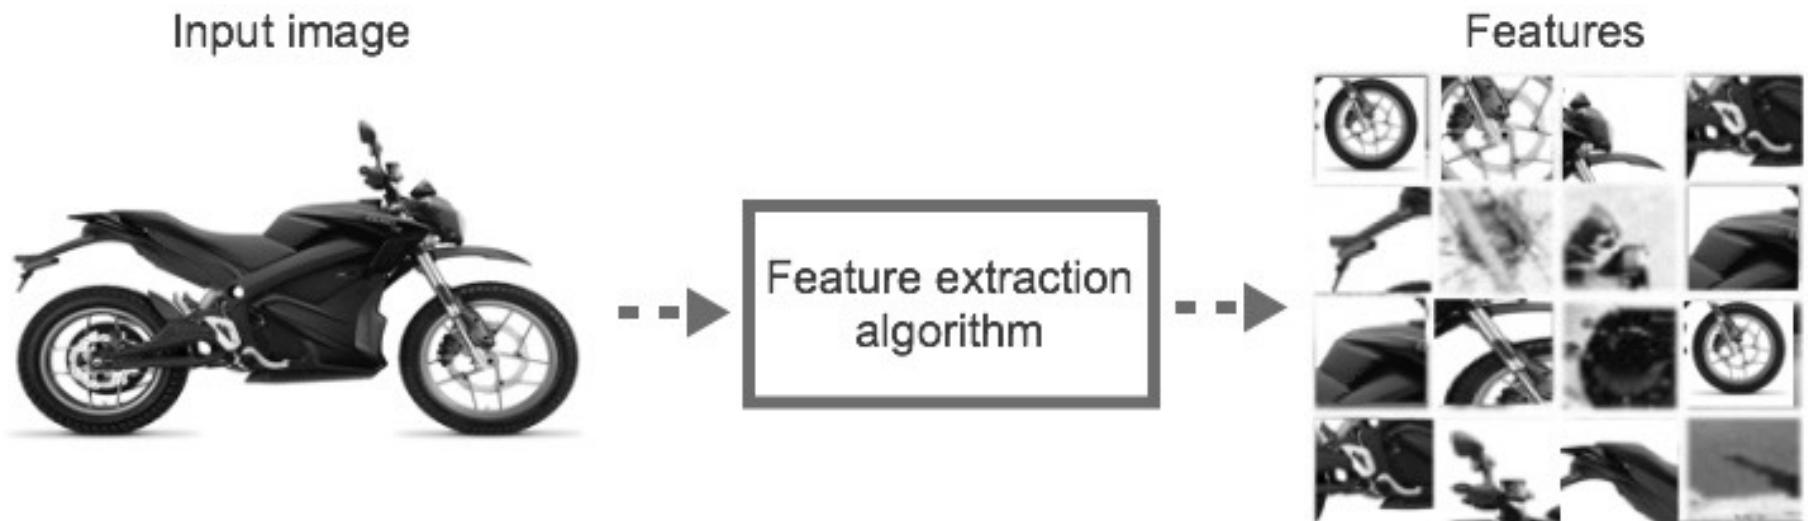
Data Cleaning

“In our experience, the tasks of exploratory data mining and data cleaning constitute **80% of the effort** that determines 80% of the value of the ultimate data.”

T. Dasu and T. Johnson
Authors of *Exploratory Data Mining and Data Cleaning*

- Good data scientists spend most of their time cleaning and formatting data.
- The rest spend most of their time complaining there is no data available.
- Data munging or data wrangling is the art of acquiring data and preparing (cleaning) it for analysis.

feature is a measurable piece of data in your image that is unique to this specific object. It maybe a distinct color in an image or a specific shape such as a line, edge, or an image segment.



features vector that is a array that makes a robust representation of the object.

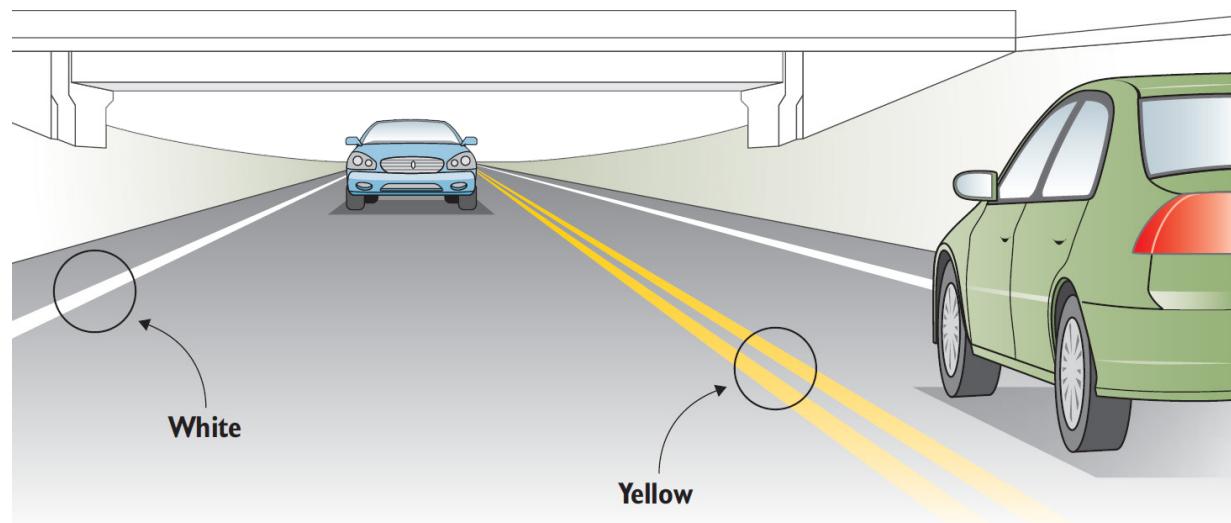
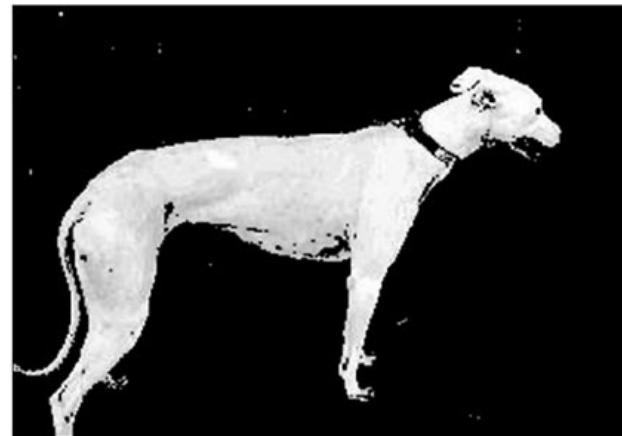
Speciailly Image

Local features

such as edge and keypoints.

Global features

such as color histogram and pixel count. By global, it is meant that the feature describes the entire image.

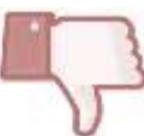


Raw data

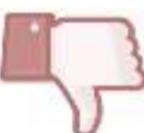
No information loss
→ Discover Unrecognized patterns



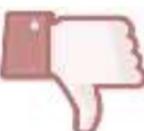
Curse of dimensionality



Higher computational cost
→ Slower learning time

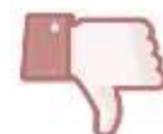


Large-scale storage is required



Featured Data

Some information loss
→ Degrade performance sometime



Dimensionality Reduction



Lower computation cost
→ Faster learning time



Small-scale storage is fine



images, even small ones, have a very large number of pixels, where each pixel is applied as an input to the model. For a grayscale image of size 100×100 pixels, there are $100 \times 100 = 10,000$ input variables to be applied to the model. For a small dataset of 100 samples, there will be a total of $100 \times 10,000 = 1,000,000$ inputs across the entire dataset. If the image is Red-Green-Blue (RGB), the total number is multiplied by 3. This requires a large memory in addition to being computationally intensive.

Another reason why feature extraction is preferred before training is that the input image has different types of objects with different properties, and **we just want to target a single object.**

repeatability

Feature after looking
at one image

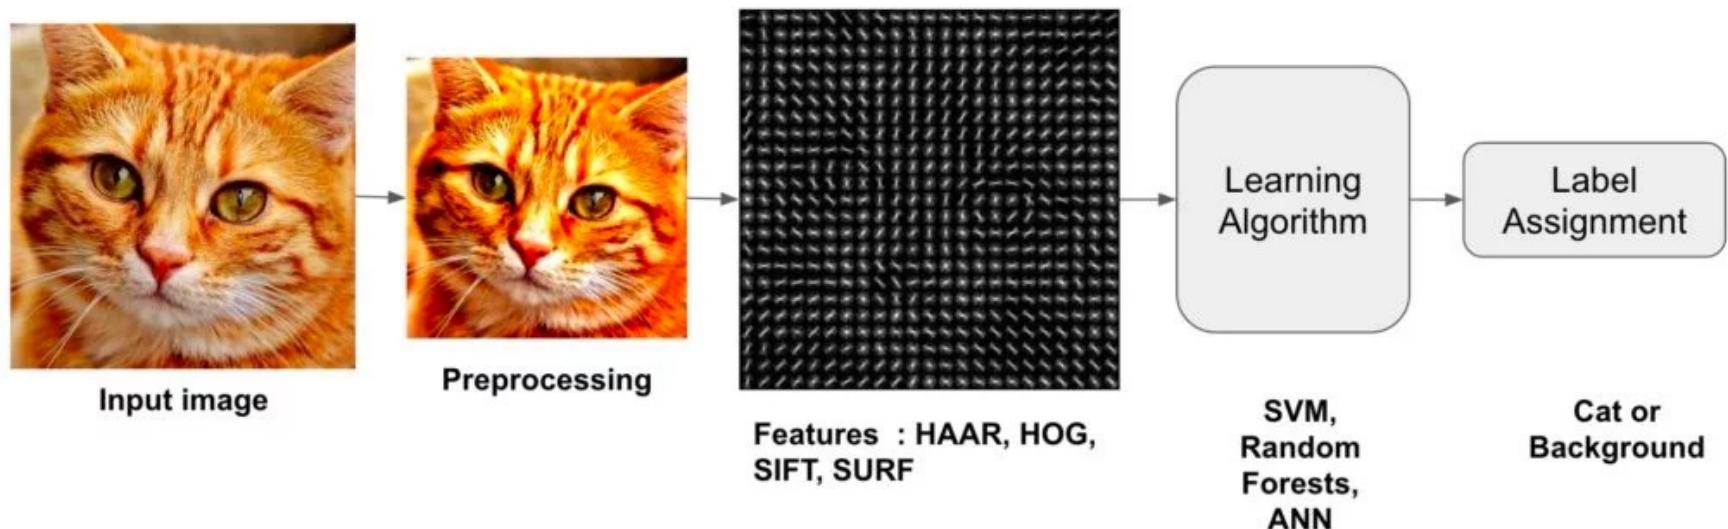


Feature after looking
at 1000s of images



A good feature will help us recognize an object in all the ways it may appear. Characteristics of a good feature:

- Identifiable
- Easily tracked and compared
- Consistent across different scales, lighting conditions, and viewing angles
- Still visible in noisy images or when only part of an object is visible



- Histogram of Oriented Gradients (HOG)
- Haar Cascades
- Scale-Invariant Feature Transform (SIFT)
- Speeded Up Robust Feature (SURF)
- Color Histogram
- Edge
- Texture
 - GLCM
 - GLGCM
 - LBP

Correlation

measurement of the similarity between two signals/sequences.
between two signals

Convolution

measurement of effect of one signal on the other signal.
between signal and a system(filter)

Theoretically, convolution are linear operations on the signal or signal modifiers, whereas correlation is a measure of similarity between two signals. As you rightly mentioned, the basic difference between convolution and correlation is that the convolution process rotates the matrix by 180 degrees..

Translational Invariance

Our vision system should be to sense, respond or detect the same object regardless of where it appears in the image.

Locality

Our vision system focusses on the local regions, without regard to what else is happening in other parts of the image.

Manul Feature Extraction

Domain knowledge

Define features

Detect features
to classify

Viewpoint variation



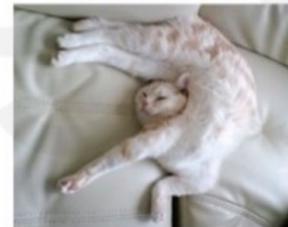
Illumination conditions



Scale variation



Deformation



Background clutter



Occlusion



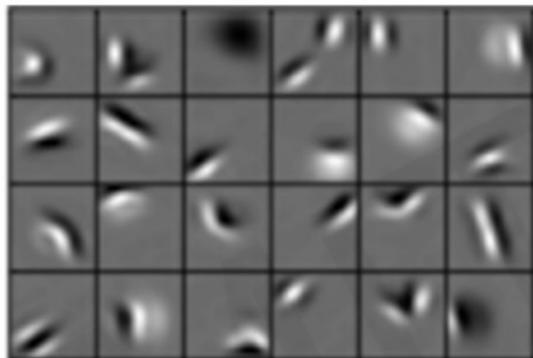
Intra-class variation



Learning Feature Representation

Can we learn a **hierarchy of features** directly from the data instead of hand engineering?

Low level features



Edges, dark spots

Mid level features



Eyes, ears, nose

High level features



Facial structure

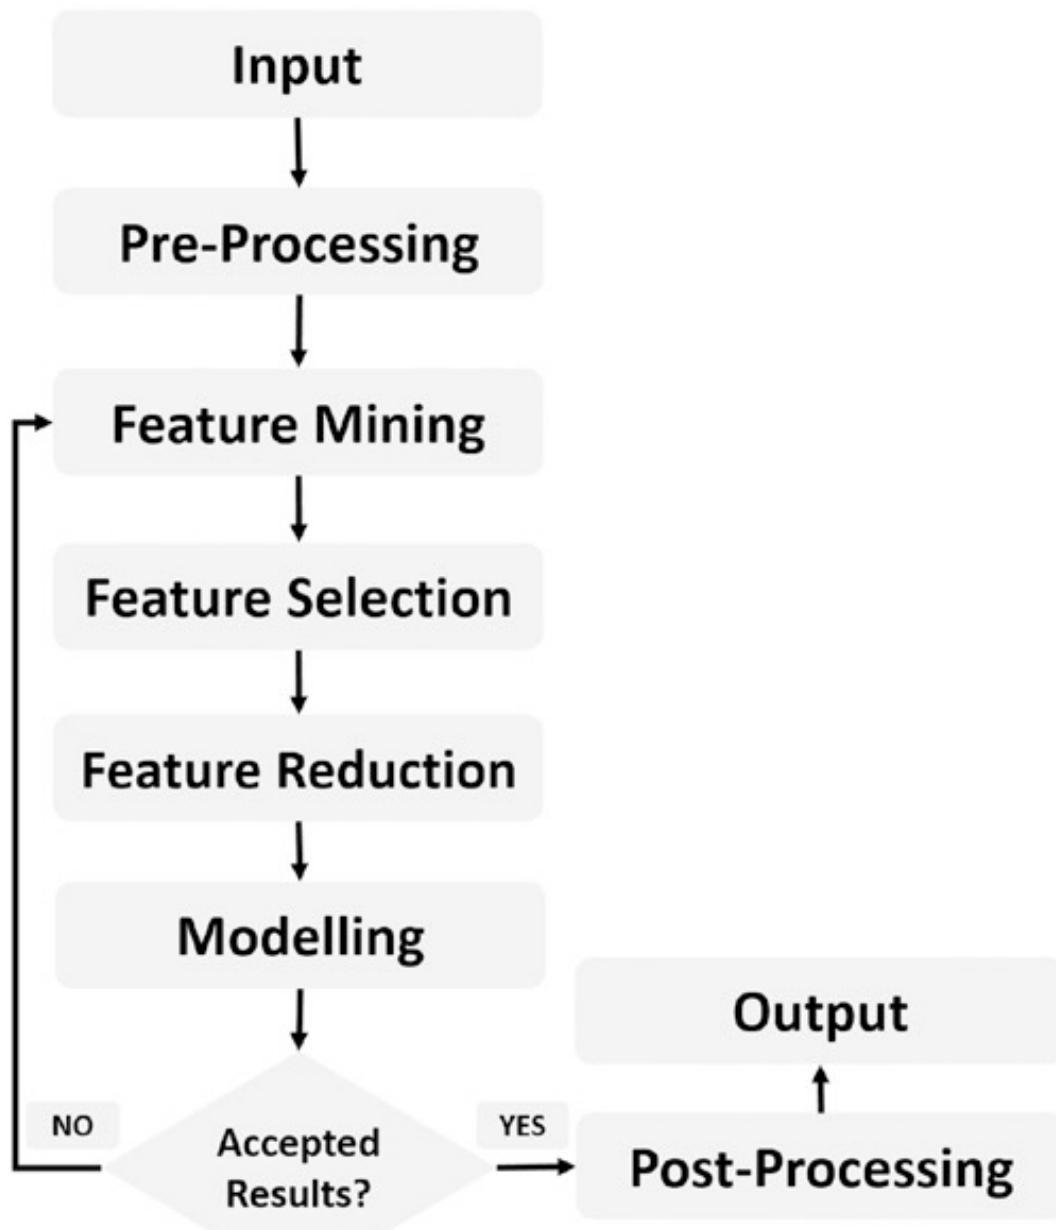
Why Deep Learning?

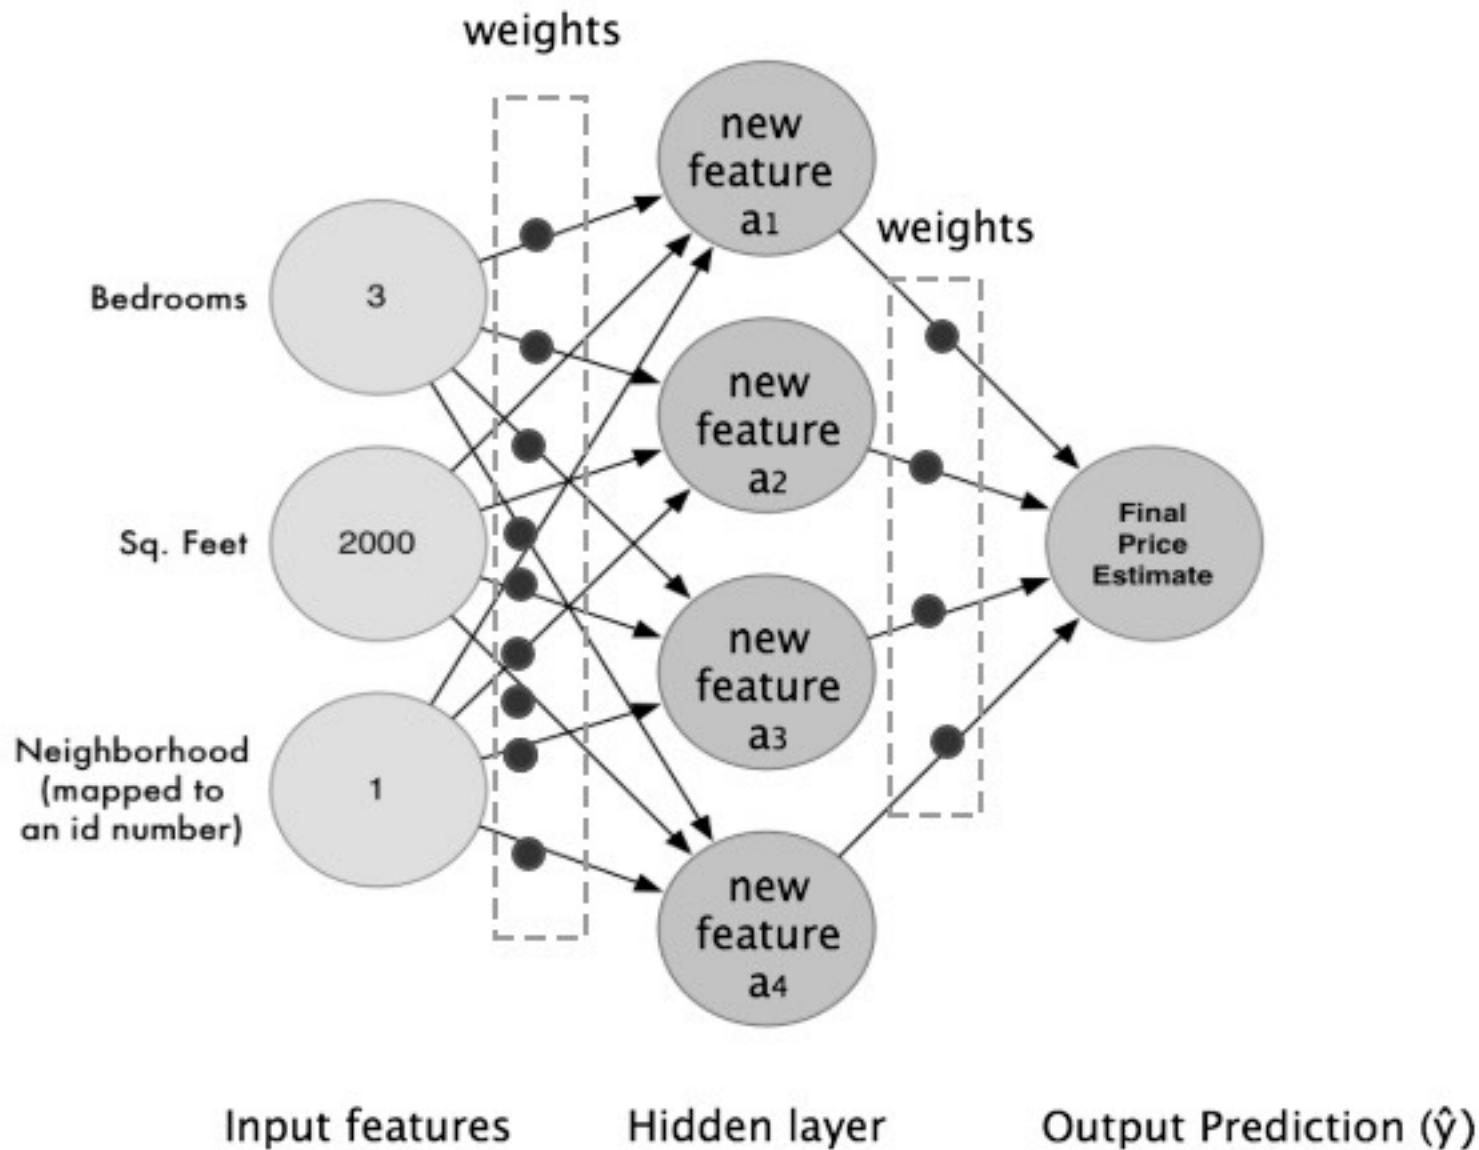
Beginning Application Development with TensorFlow and Keras

3-5

<https://brunch.co.kr/@itschloe1/8>

<http://hyperparameter.space/blog/when-not-to-use-deep-learning/>

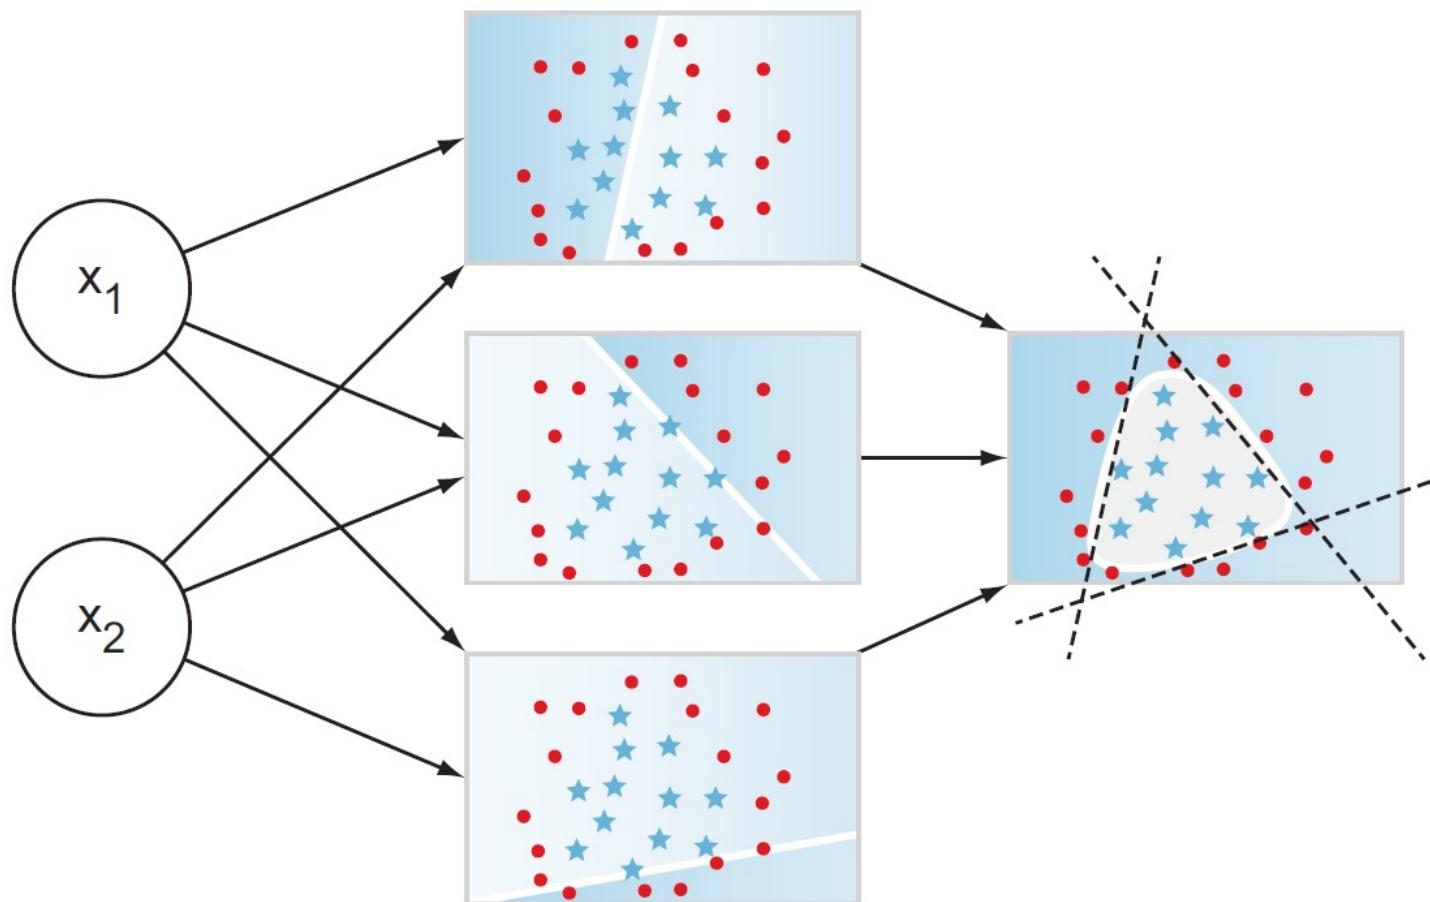




Input features

Hidden layer

Output



Feature Crosses

A **feature cross** is a **synthetic feature** formed by multiplying (crossing) two or more features. Crossing combinations of features can provide predictive abilities beyond what those features can provide individually.

Kinds of feature crosses

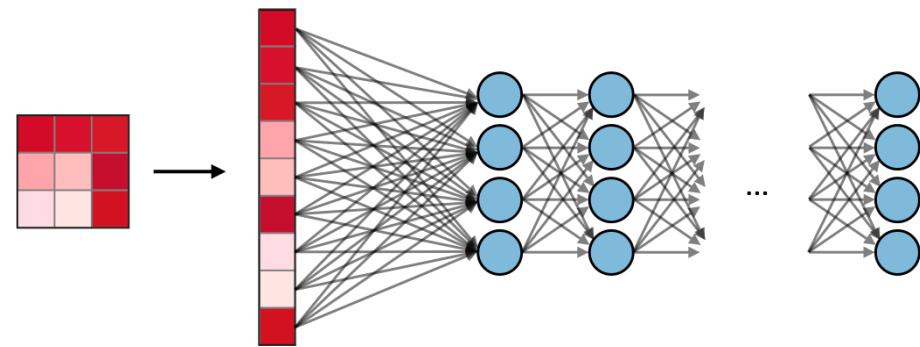
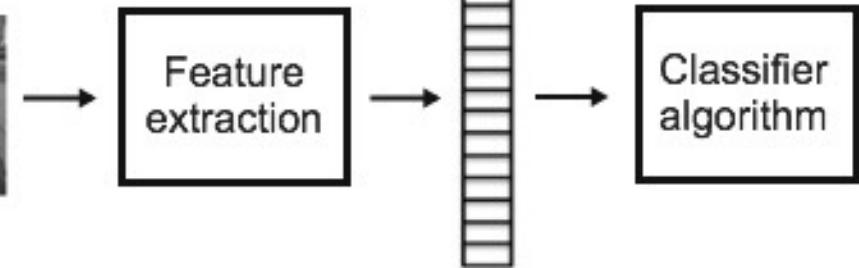
We can create many different kinds of feature crosses. For example:

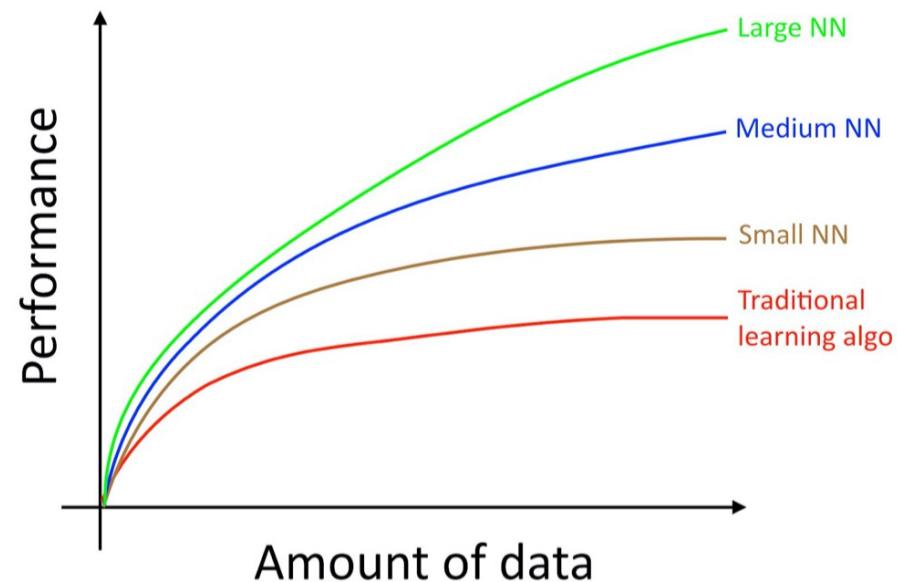
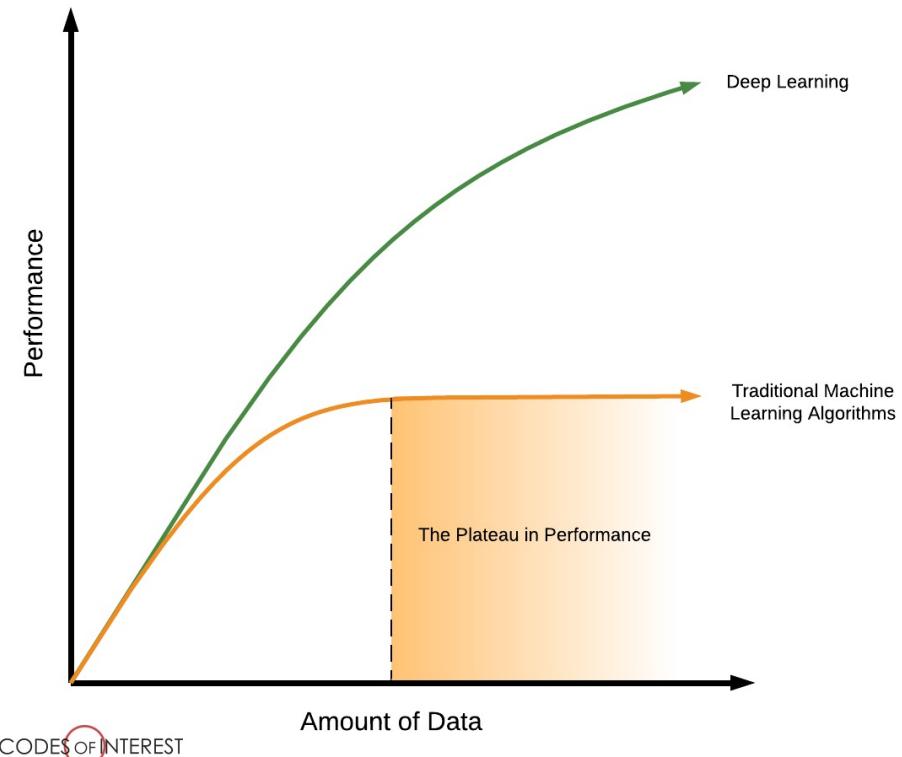
- [A X B]: a feature cross formed by multiplying the values of two features.
- [A x B x C x D x E]: a feature cross formed by multiplying the values of five features.
- [A x A]: a feature cross formed by squaring a single feature.

Images dataset of 10,000 images



1D Features vector

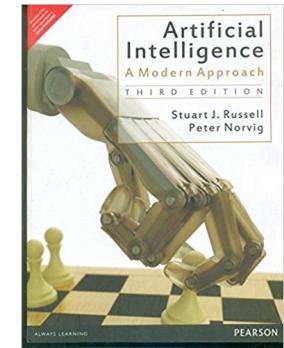




Big Data

The Unreasonable Effectiveness of Data (2009)

Peter Norvig



Learning Curves for Confusion Set Disambiguation (2001)

Michele Banko & Eric Brill

Big data

No overfitting to training data



Working with complex models for difficult problems

Higher computational cost
→ Slower learning time

Large-scale storage is required

Small Data

Overfitting to training data is possible



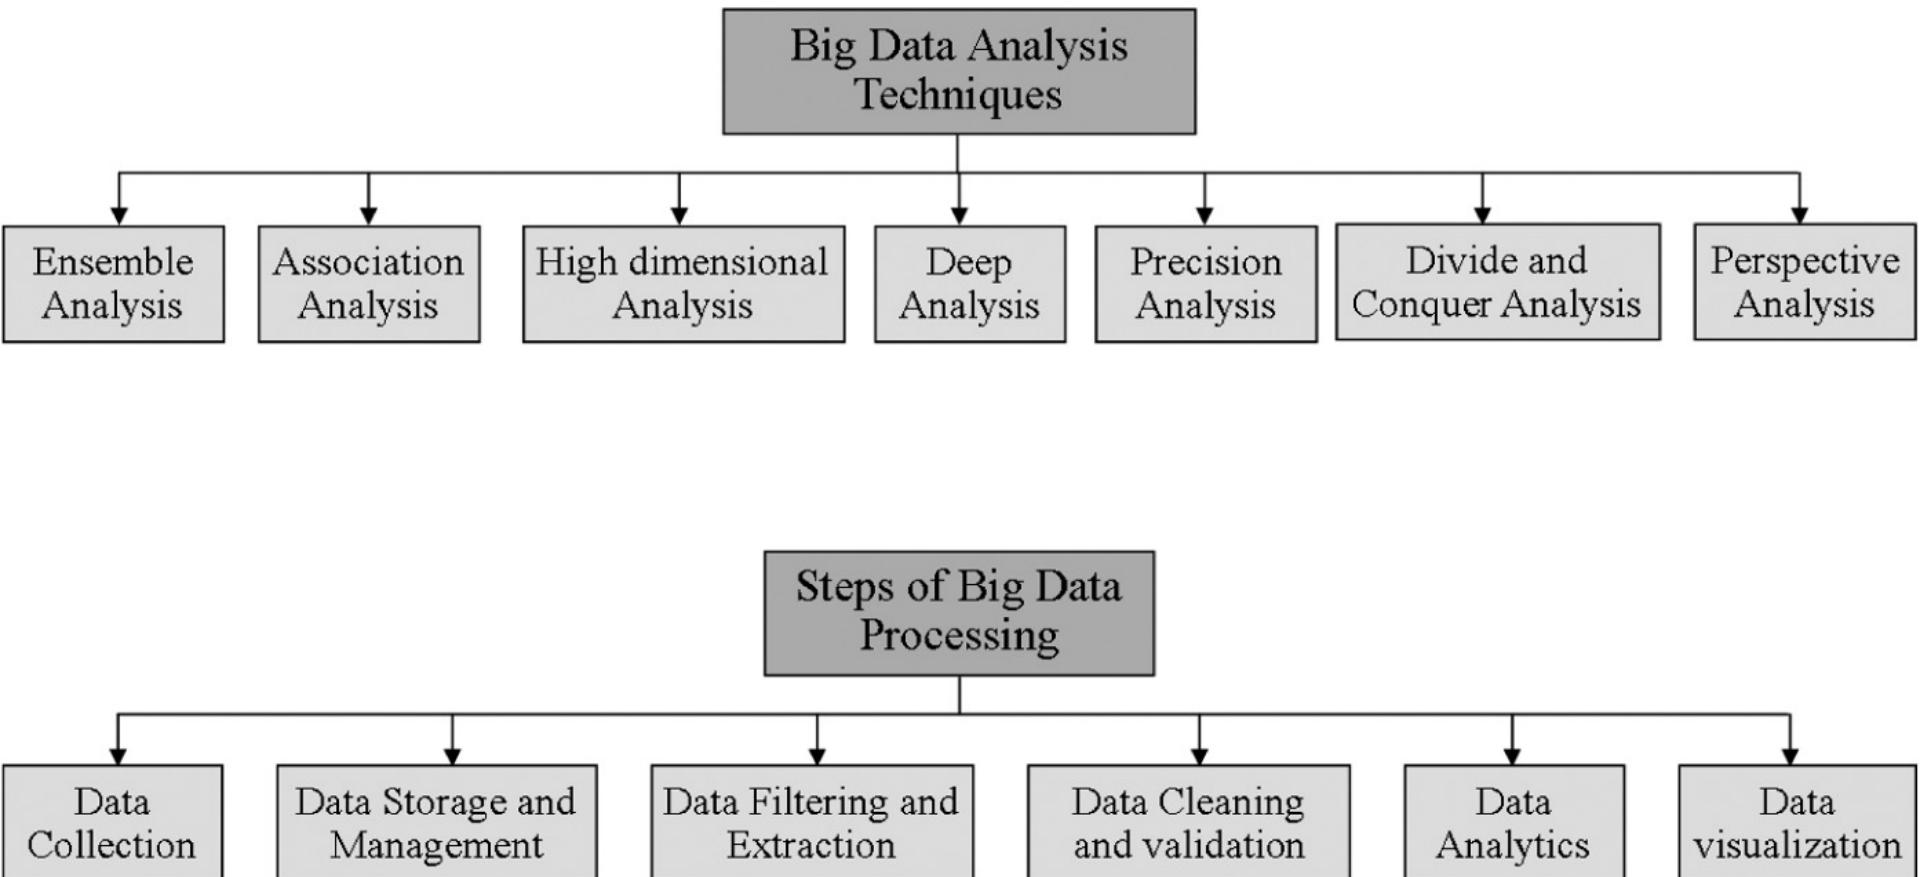
Working with simple models for easy problems

Lower computation cost
→ Faster learning time

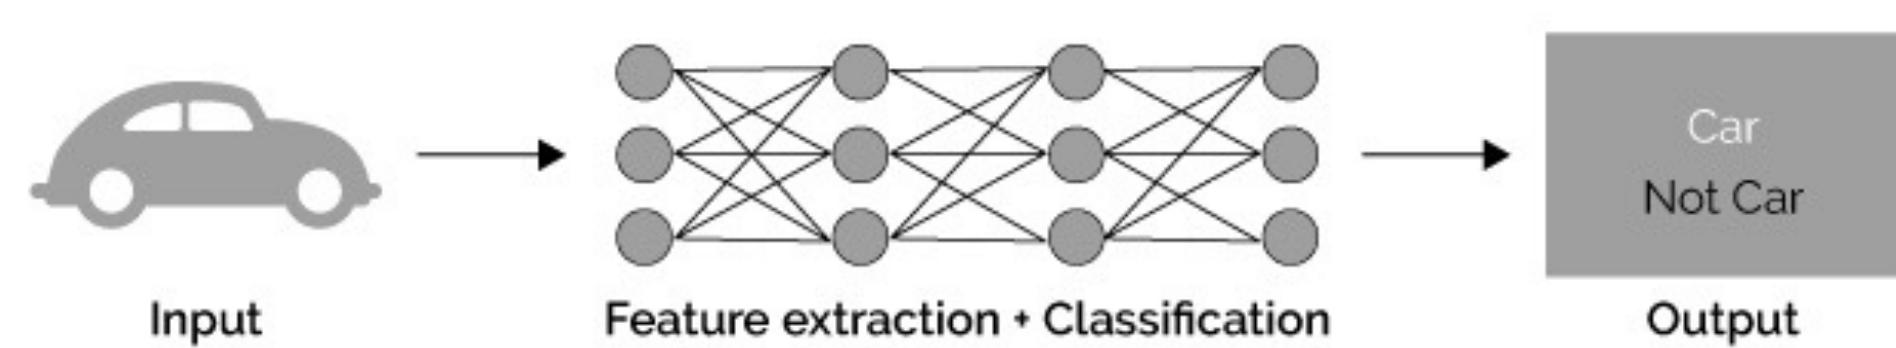


Small-scale storage is fine

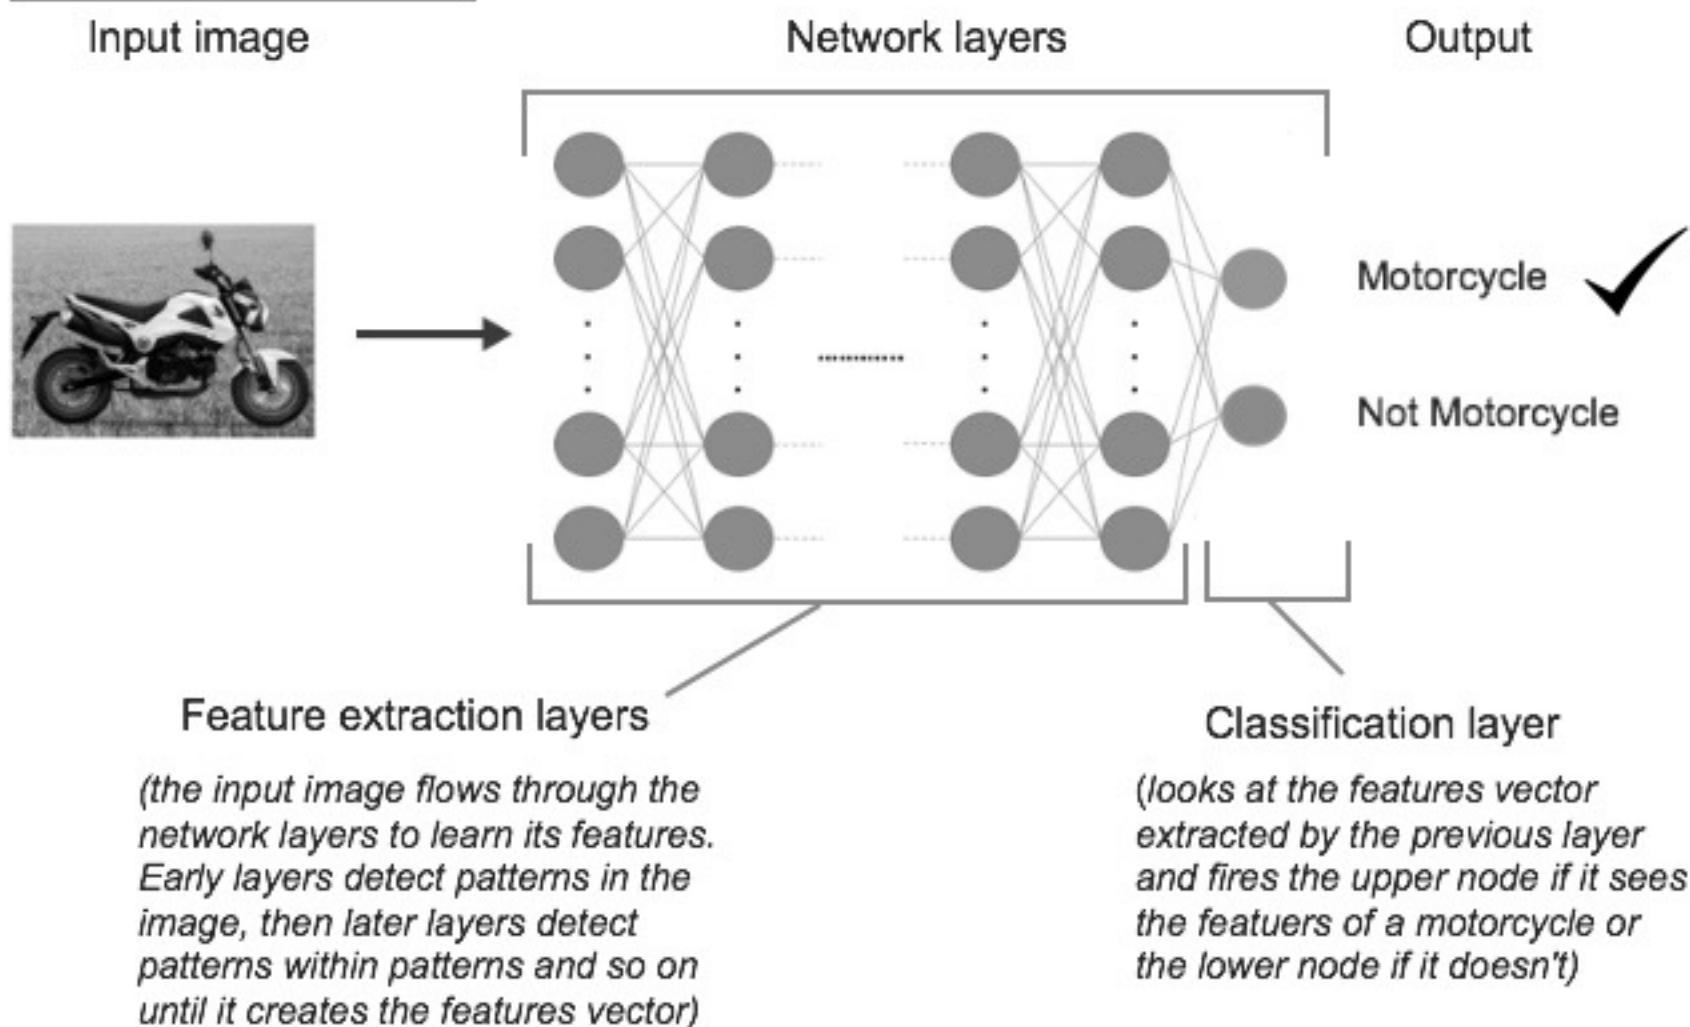


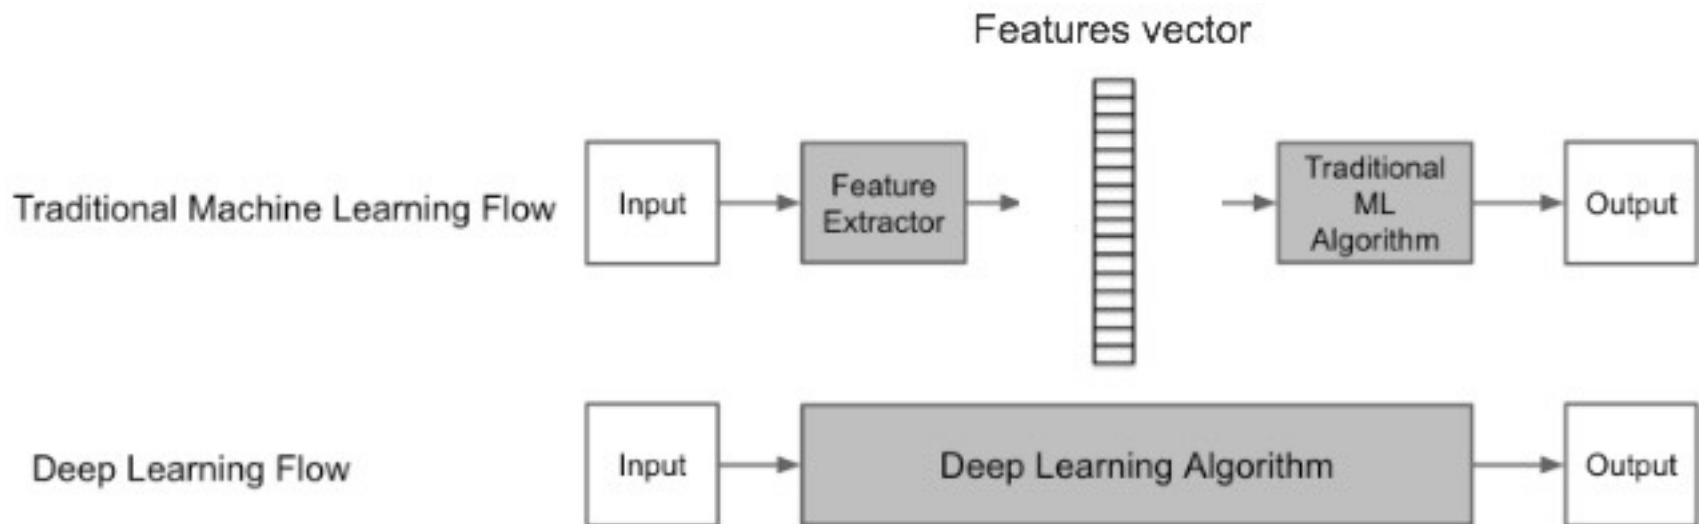


- The Promise of Automatic Feature Extraction. Features can be automatically learned and extracted from raw image data.
- The Promise of End-to-End Models. Single end-to-end models can replace pipelines of specialized models.
- The Promise of Model Reuse. Learned features and even entire models can be reused across tasks.
- The Promise of Superior Performance. Techniques demonstrate better skill on challenging tasks.
- The Promise of General Method. A single general method can be used on a range of related tasks.



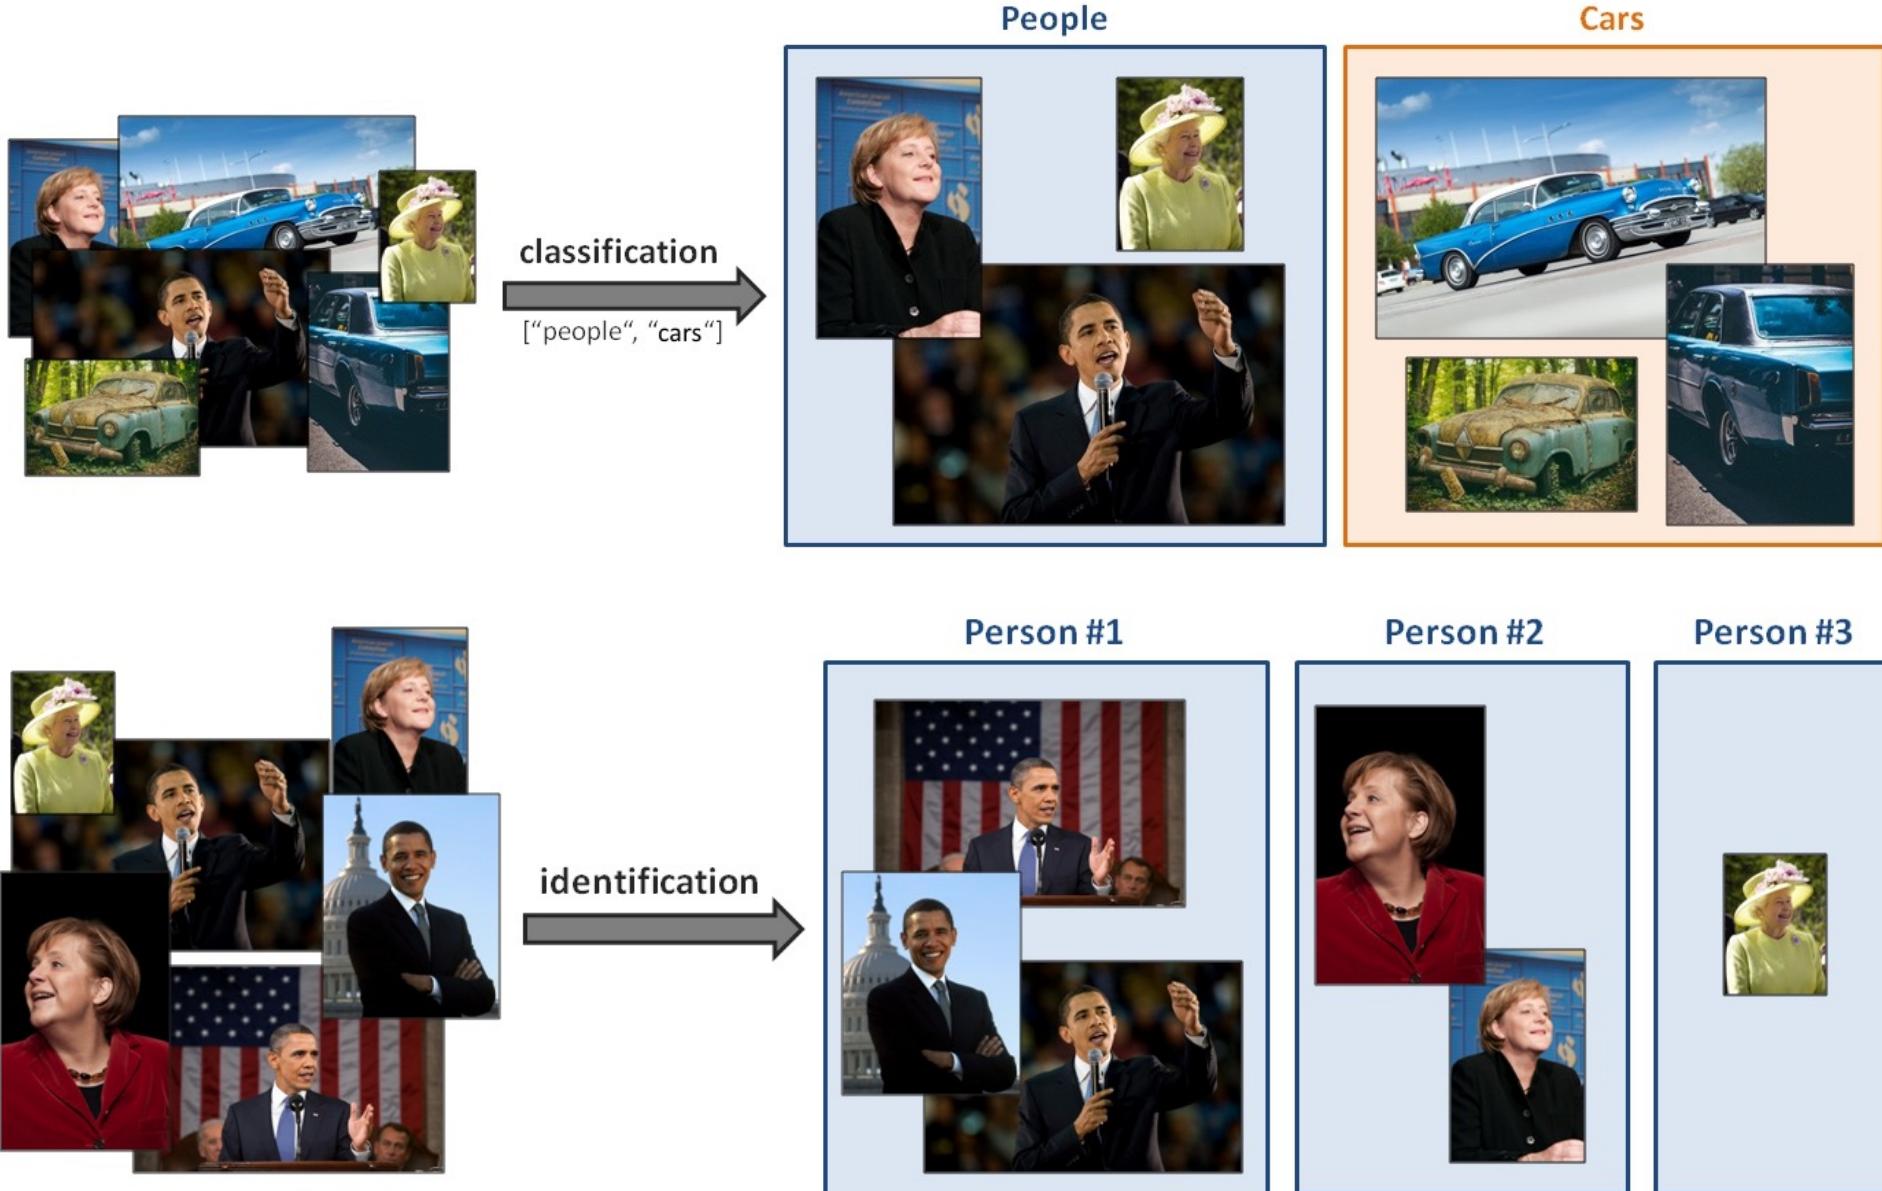
Deep Learning Classifier





CV Application

- Object Classification: What broad category of object is in this photograph?
- Object Identification: Which type of a given object is in this photograph?
- Object Verification: Is the object in the photograph?
- Object Detection: Where are the objects in the photograph?
- Object Landmark Detection: What are the key points for the object in the photograph?
- Object Segmentation: What pixels belong to the object in the image?
- Object Recognition: What objects are in this photograph and where are they?



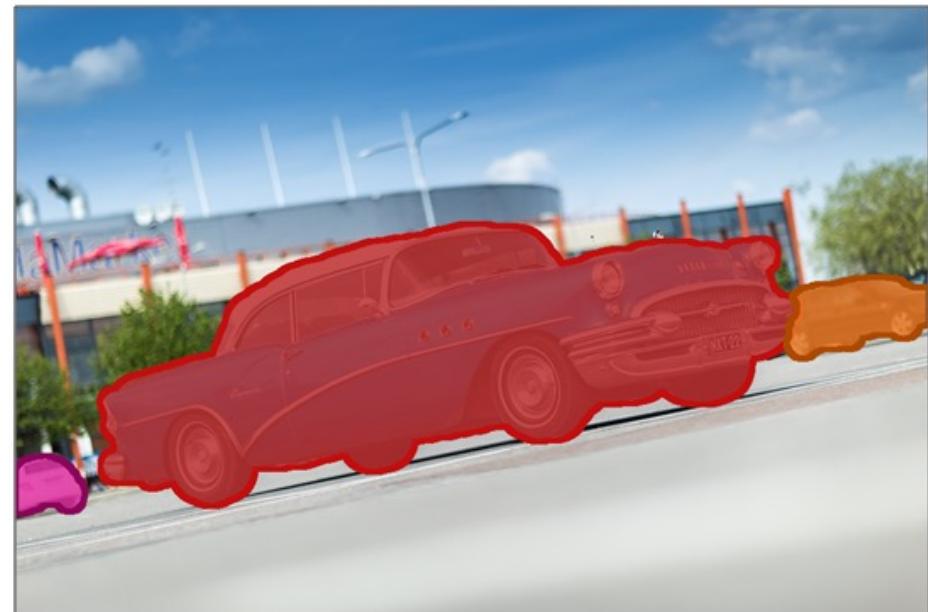
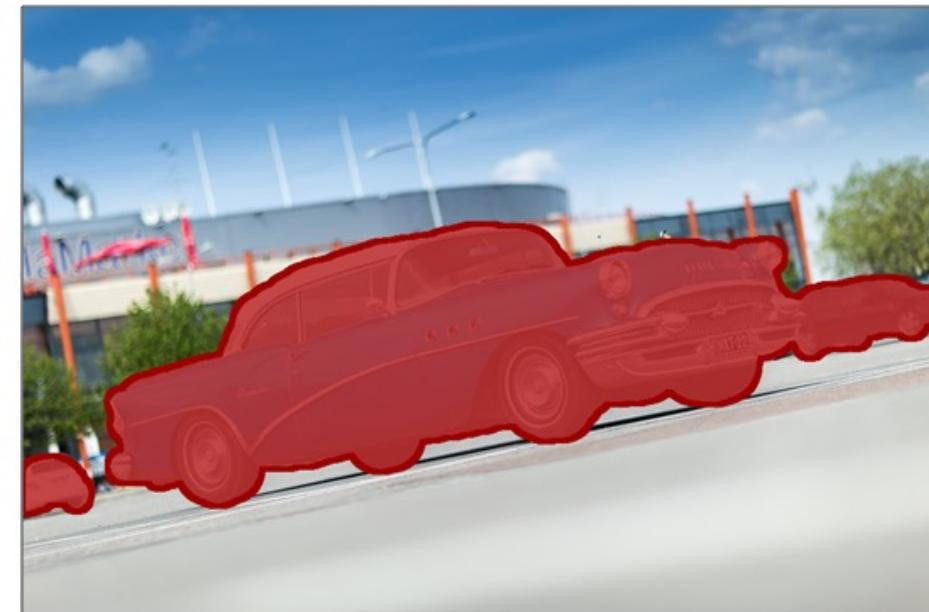


detection
“car”



object segmentation

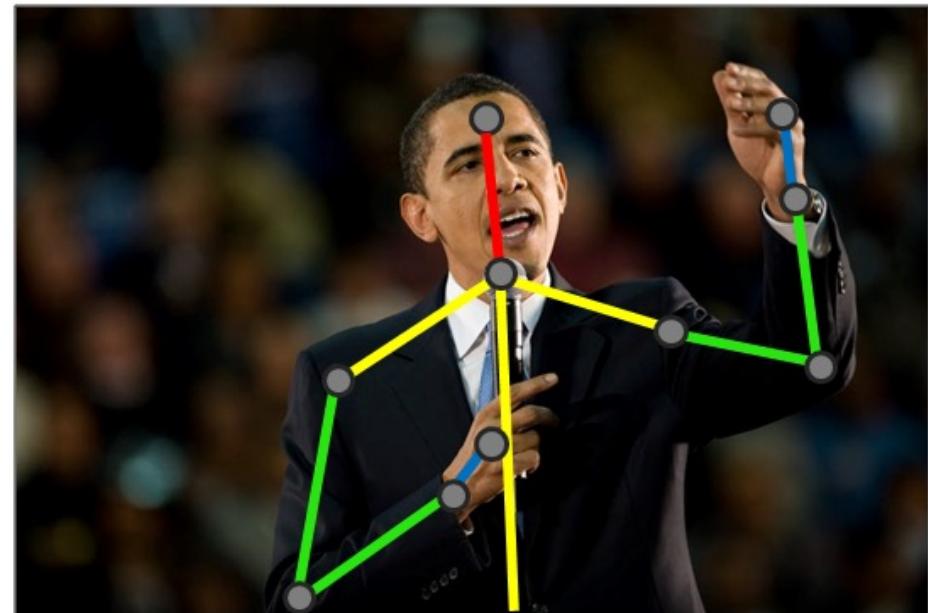
instance segmentation



rigid object pose estimation



human pose estimation





A person riding a motorcycle on a dirt road.



A group of young people playing a game of frisbee.

Original image



+

Style



=

Generated art



This small blue bird has a short pointy beak and brown on its wings



This bird is completely red with black wings and pointy beak

